

FINAL-PROJECT

GROUP - 9

2023-05-02

Importing of required libraries for the visualization and interpretation.

```
library(dplyr)
library(tidyr)
library(tidyverse)
library(leaps)
library(knitr)
library(ggplot2)
library("reshape2")
library(caret)
library(psych)
library(tree)
library(rpart)
library(rattle)
library(randomForest)
library(lattice)
library(plotly)
library(corrplot)
library(car)
library(skimr)
library(ggsci)
library(moments)
library(randomForestSRC)
library(gridExtra)
library(class)
library(car)
```

Importing the dataset from the current directory.

```
Heart_disease = read.csv("/Users/vishaypaka/Documents/STAT-515/Datasets/heart_failure_clinical_records_0
```

PREPROCESSING OF THE DATASET

```
table(is.na(Heart_disease)) # no null values
```

```
##
## FALSE
## 3887
```

We can see that our dataset doesn't contain any null values. Now let's observe what datatype are our predictor variables and if there are any categorical or binary variables, let's factor them.

Structure of the dataset

```
#Exploring the dataset
str(Heart_disease) #checking the datatypes of variables present
```

```
## 'data.frame':    299 obs. of  13 variables:
## $ age           : num  75 55 65 50 65 90 75 60 65 80 ...
## $ anaemia       : int   0 0 0 1 1 1 1 0 1 ...
## $ creatinine_phosphokinase: int 582 7861 146 111 160 47 246 315 157 123 ...
## $ diabetes      : int   0 0 0 0 1 0 0 1 0 0 ...
## $ ejection_fraction : int 20 38 20 20 20 40 15 60 65 35 ...
## $ high_blood_pressure : int 1 0 0 0 0 1 0 0 0 1 ...
## $ platelets      : num 265000 263358 162000 210000 327000 ...
## $ serum_creatinine : num 1.9 1.1 1.3 1.9 2.7 2.1 1.2 1.1 1.5 9.4 ...
## $ serum_sodium    : int 130 136 129 137 116 132 137 131 138 133 ...
## $ sex             : int 1 1 1 1 0 1 1 1 0 1 ...
## $ smoking         : int 0 0 1 0 0 1 0 1 0 1 ...
## $ time            : int 4 6 7 7 8 8 10 10 10 10 ...
## $ DEATH_EVENT     : int 1 1 1 1 1 1 1 1 1 1 ...
```

Data pre-processing

```
table(Heart_disease$platelets)
```

```
##
##      25100      47000      51000      62000      70000      73000      75000      87000
##          1          1          1          1          1          1          1          1
##     105000     119000     122000     126000     127000     130000     132000     133000
##          1          1          1          1          2          1          1          2
##     136000     140000     141000     147000     149000     150000     151000     153000
##          1          2          1          2          3          1          1          3
##     155000     160000     162000     164000     166000     172000     173000     174000
##          1          1          2          1          2          2          2          1
##     176000     179000     181000     184000     185000     186000     188000     189000
##          1          1          1          1          2          1          1          3
##     192000     194000     196000     198000     2e+05     201000     203000     204000
##          1          3          2          1          1          1          3          2
##     208000     210000     211000     212000     213000     215000     216000     217000
##          1          3          1          1          1          2          2          1
##     218000     219000     220000     221000     222000     223000     224000     225000
##          2          2          3          4          2          3          1          1
##     226000     227000     228000     229000     231000     232000     233000     235000
##          4          1          4          1          2          1          1          4
##     236000     237000     241000     242000     243000     244000     246000     248000
##          1          4          1          2          1          3          1          1
##     249000     250000     252000     253000     254000     255000     257000     259000
##          3          1          1          2          3          4          1          1
```

##	260000	262000	263000	263358.03	264000	265000	266000	267000
##	1	2	1	25	1	3	2	2
##	268000	270000	271000	274000	275000	276000	277000	279000
##	1	2	4	3	1	2	2	4
##	281000	282000	283000	284000	286000	289000	290000	293000
##	1	1	3	1	1	1	1	1
##	294000	295000	297000	298000	3e+05	301000	302000	303000
##	1	1	2	1	1	1	3	1
##	304000	305000	306000	308000	309000	310000	314000	317000
##	2	4	1	1	1	1	1	1
##	318000	319000	321000	324000	325000	327000	328000	329000
##	1	2	1	1	1	3	1	2
##	330000	334000	336000	337000	338000	348000	350000	351000
##	1	2	1	1	1	1	1	2
##	358000	360000	362000	365000	368000	371000	374000	377000
##	1	1	3	2	2	1	1	1
##	382000	385000	388000	389000	390000	395000	404000	406000
##	1	1	1	2	2	2	1	2
##	418000	422000	427000	448000	451000	454000	461000	481000
##	1	1	1	1	2	1	1	1
##	497000	504000	507000	533000	543000	621000	742000	850000
##	1	1	1	1	1	1	1	1

```
Heart_disease$platelets <- round(Heart_disease$platelets)
```

```
#Formatting the scientific notations for the platelets variable.
```

```
Heart_disease$platelets <- format(Heart_disease$platelets, scientific = FALSE)
```

```
table(Heart_disease$platelets)
```

##	25100	47000	51000	62000	70000	73000	75000	87000	105000	119000	122000
##	1	1	1	1	1	1	1	1	1	1	1
##	126000	127000	130000	132000	133000	136000	140000	141000	147000	149000	150000
##	1	2	1	1	2	1	2	1	2	3	1
##	151000	153000	155000	160000	162000	164000	166000	172000	173000	174000	176000
##	1	3	1	1	2	1	2	2	2	1	1
##	179000	181000	184000	185000	186000	188000	189000	192000	194000	196000	198000
##	1	1	1	2	1	1	3	1	3	2	1
##	200000	201000	203000	204000	208000	210000	211000	212000	213000	215000	216000
##	1	1	3	2	1	3	1	1	1	2	2
##	217000	218000	219000	220000	221000	222000	223000	224000	225000	226000	227000
##	1	2	2	3	4	2	3	1	1	4	1
##	228000	229000	231000	232000	233000	235000	236000	237000	241000	242000	243000
##	4	1	2	1	1	4	1	4	1	2	1
##	244000	246000	248000	249000	250000	252000	253000	254000	255000	257000	259000
##	3	1	1	3	1	1	2	3	4	1	1
##	260000	262000	263000	263358	264000	265000	266000	267000	268000	270000	271000
##	1	2	1	25	1	3	2	2	1	2	4
##	274000	275000	276000	277000	279000	281000	282000	283000	284000	286000	289000
##	3	1	2	2	4	1	1	3	1	1	1
##	290000	293000	294000	295000	297000	298000	300000	301000	302000	303000	304000
##	1	1	1	1	2	1	1	1	3	1	2
##	305000	306000	308000	309000	310000	314000	317000	318000	319000	321000	324000

```
##      4      1      1      1      1      1      1      1      2      1      1
## 325000 327000 328000 329000 330000 334000 336000 337000 338000 348000 350000
##      1      3      1      2      1      2      1      1      1      1      1
## 351000 358000 360000 362000 365000 368000 371000 374000 377000 382000 385000
##      2      1      1      3      2      2      1      1      1      1      1
## 388000 389000 390000 395000 404000 406000 418000 422000 427000 448000 451000
##      1      2      2      2      1      2      1      1      1      1      2
## 454000 461000 481000 497000 504000 507000 533000 543000 621000 742000 850000
##      1      1      1      1      1      1      1      1      1      1      1
```

```
#rounding off the value of an age variable.
table(Heart_disease$age)
```

```
##
##      40      41      42      43      44      45      46      47      48      49      50
##       7       1       7       1       2      19       3       1       2       4      27
##      51      52      53      54      55      56      57      58      59      60 60.667
##       4       5      10       2      17       1       2      10       4      33       2
##      61      62      63      64      65      66      67      68      69      70      72
##       4       5       8       3      26       2       2       5       3      25       7
##      73      75      77      78      79      80      81      82      85      86      87
##       4      11       2       2       1       7       1       3       6       1       1
##      90      94      95
##       3       1       2
```

```
Heart_disease$age <- round(Heart_disease$age)

table(Heart_disease$age)
```

```
##
## 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65
##  7  1  7  1  2 19  3  1  2  4 27  4  5 10  2 17  1  2 10  4 33  6  5  8  3 26
## 66 67 68 69 70 72 73 75 77 78 79 80 81 82 85 86 87 90 94 95
##  2  2  5  3 25  7  4 11  2  2  1  7  1  3  6  1  1  3  1  2
```

Here when we observe the column of age one observation as age 66.67, which stands out of the rest, so rounding off as we have less observations for prediction. Also platelets column has scientific notation for two records. So, formatting it into numeric as an pre-processing step.

Let's factorize the binary variables and convert continuous to numeric.

```
Heart_disease$anaemia = as.factor(Heart_disease$anaemia)
Heart_disease$diabetes = as.factor(Heart_disease$diabetes)
Heart_disease$high_blood_pressure = as.factor(Heart_disease$high_blood_pressure)
Heart_disease$sex = as.factor(Heart_disease$sex)
Heart_disease$smoking = as.factor(Heart_disease$smoking)
Heart_disease$DEATH_EVENT = as.factor(Heart_disease$DEATH_EVENT)
Heart_disease$creatinine_phosphokinase = as.numeric(Heart_disease$creatinine_phosphokinase)
Heart_disease$ejection_fraction = as.numeric(Heart_disease$ejection_fraction)
Heart_disease$serum_sodium = as.numeric(Heart_disease$serum_sodium)
Heart_disease$platelets = as.numeric(Heart_disease$platelets)
str(Heart_disease)
```

```
## 'data.frame': 299 obs. of 13 variables:
## $ age : num 75 55 65 50 65 90 75 60 65 80 ...
## $ anaemia : Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 2 1 2 ...
## $ creatinine_phosphokinase: num 582 7861 146 111 160 ...
## $ diabetes : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 2 1 1 ...
## $ ejection_fraction : num 20 38 20 20 20 40 15 60 65 35 ...
## $ high_blood_pressure : Factor w/ 2 levels "0","1": 2 1 1 1 1 2 1 1 1 2 ...
## $ platelets : num 265000 263358 162000 210000 327000 ...
## $ serum_creatinine : num 1.9 1.1 1.3 1.9 2.7 2.1 1.2 1.1 1.5 9.4 ...
## $ serum_sodium : num 20 38 20 20 20 40 15 60 65 35 ...
## $ sex : Factor w/ 2 levels "0","1": 2 2 2 2 1 2 2 2 1 2 ...
## $ smoking : Factor w/ 2 levels "0","1": 1 1 2 1 1 2 1 2 1 2 ...
## $ time : int 4 6 7 7 8 8 10 10 10 10 ...
## $ DEATH_EVENT : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
```

Producing some numerical and graphical summaries of the data set.

Statistical Analysis

```
summary(Heart_disease) #summary of the dataset
```

```
##      age      anaemia creatinine_phosphokinase diabetes ejection_fraction
## Min.   :40.00   0:170   Min.    : 23.0           0:174   Min.    :14.00
## 1st Qu.:51.00   1:129   1st Qu.: 116.5           1:125   1st Qu.:30.00
## Median :60.00           Median : 250.0           Median :38.00
## Mean   :60.84           Mean   : 581.8           Mean   :38.08
## 3rd Qu.:70.00           3rd Qu.: 582.0           3rd Qu.:45.00
## Max.   :95.00           Max.   :7861.0           Max.   :80.00
## high_blood_pressure platelets      serum_creatinine serum_sodium sex
## 0:194           Min.    : 25100   Min.    :0.500   Min.    :14.00  0:105
## 1:105           1st Qu.:212500   1st Qu.:0.900   1st Qu.:30.00  1:194
##               Median :262000   Median :1.100   Median :38.00
##               Mean   :263358   Mean   :1.394   Mean   :38.08
##               3rd Qu.:303500   3rd Qu.:1.400   3rd Qu.:45.00
##               Max.   :850000   Max.   :9.400   Max.   :80.00
## smoking      time      DEATH_EVENT
## 0:203   Min.    : 4.0   0:203
## 1: 96   1st Qu.: 73.0   1: 96
##               Median :115.0
##               Mean   :130.3
##               3rd Qu.:203.0
##               Max.   :285.0
```

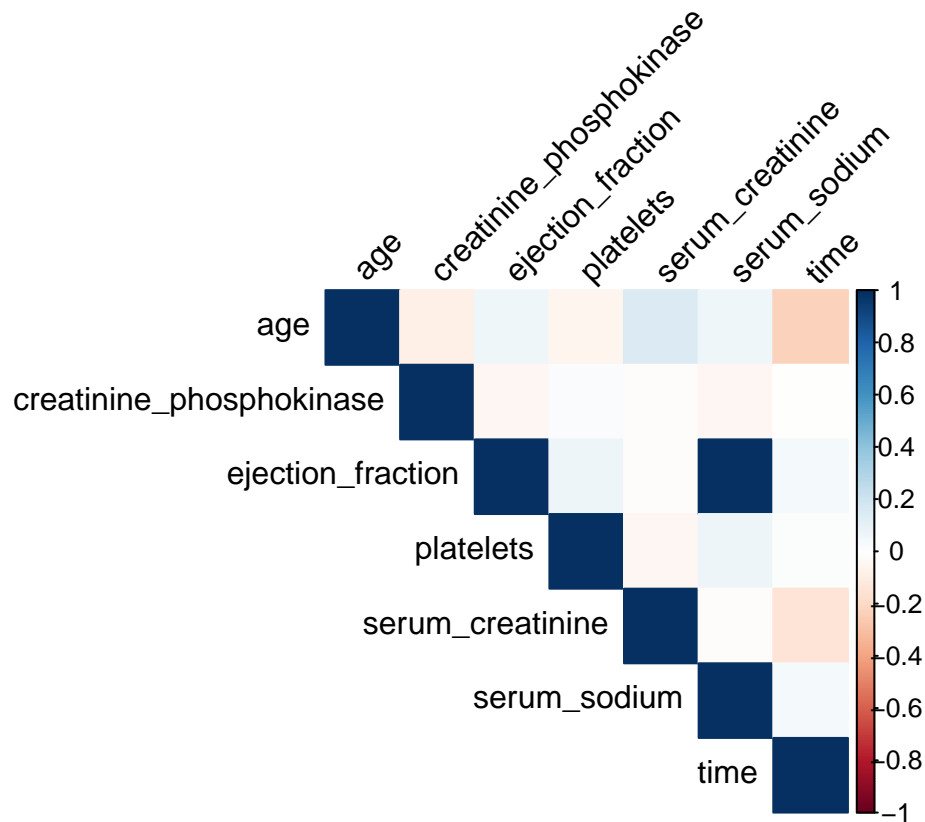
By summary of the dataset, we can observe all the variable's mean, median, min and max. We can see there are patients having anaemia are 129 and does not have anaemia are 170. 125 patients have diabetes and 174 doesn't have diabetes 194 patients dose not have high bp and 105 patients have high bp. out of all patients 203 dose not smoke and the rest of them i.e; 96 people will smoke. In our dataset we have 105 patients recorded as female and 194 as of male. The total death events taken place during the follow up period are 96 and the patients who survived are 203.

By observing all these statistical summaries of the predictors and response variables. We can conclude that the data set is slightly imbalanced as there are more observation of the patients who survived compare to the

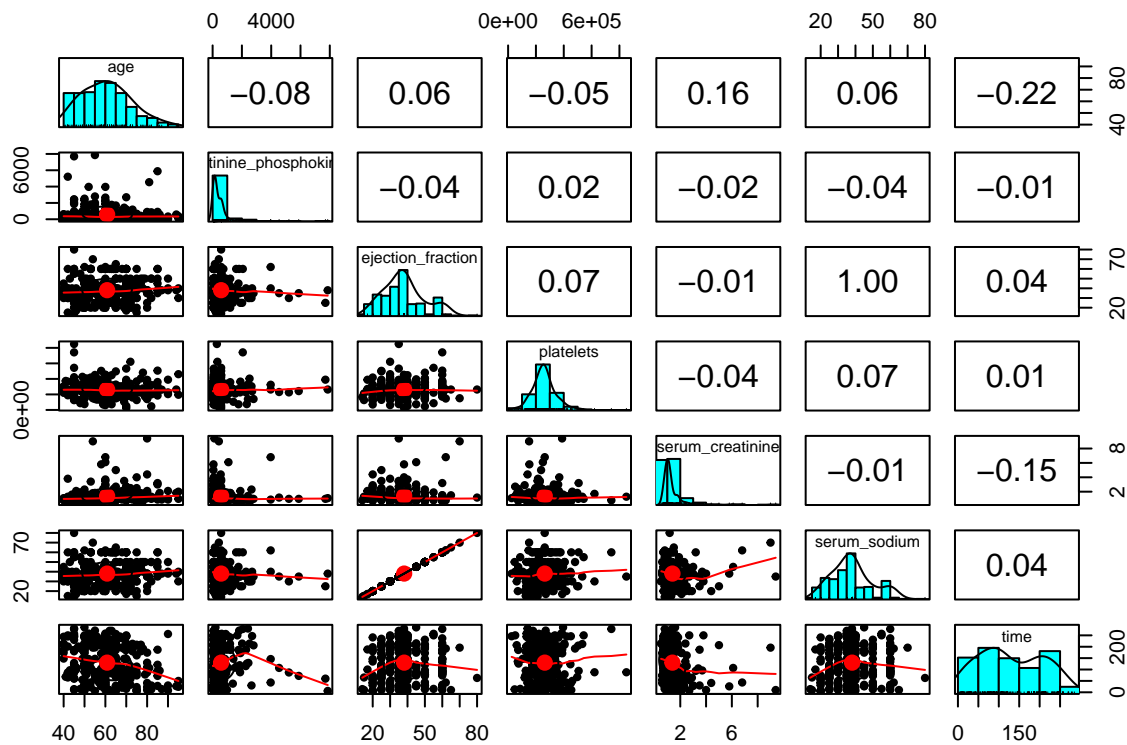
people who died because of the heart failure but have less difference of records with patients having anaemia or not, diabetes or not, high blood pressure or not, smoking or not. In general out of 7.2 billion population only less proportionate of people will die with heart failure condition.

Correlation analysis

```
numeric_cols <- Heart_disease %>% select_if(is.numeric)
cor_mat <- cor(numeric_cols)
corrplot(cor_mat, method = "color", type = "upper", tl.col = "black", tl.srt = 45)
```



```
Heart_disease1 = subset(Heart_disease, select = -c(anaemia,high_blood_pressure,sex,smoking,DEATH_EVENT,
pairs.panels(Heart_disease1)
```

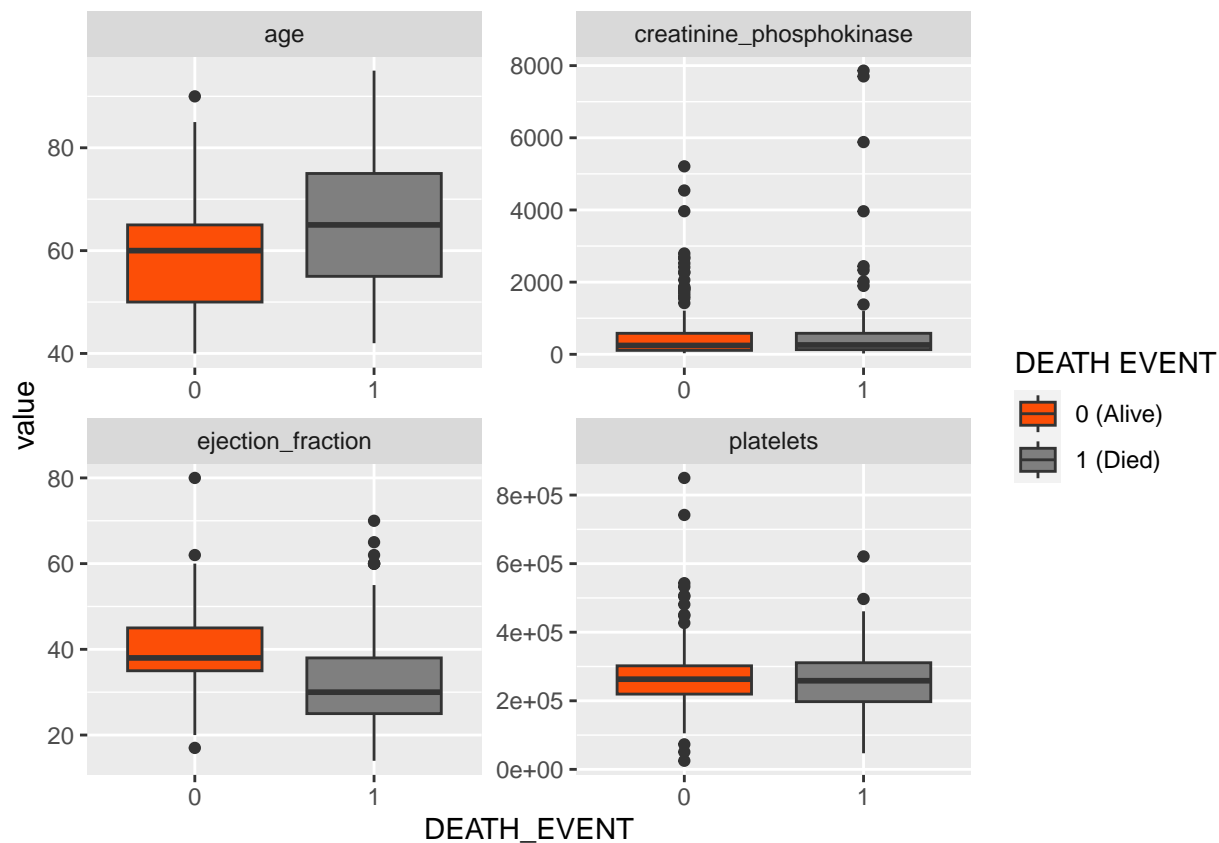


The correlation between serum sodium and ejection fraction is 1.0, which indicates that they are perfectly linearly related and increase or decrease together. Having correlation 1 means they are perfect positive and from the chart above we can see a linear line between serum sodium and ejection fraction. Let's see how it affects our models when fitted later.

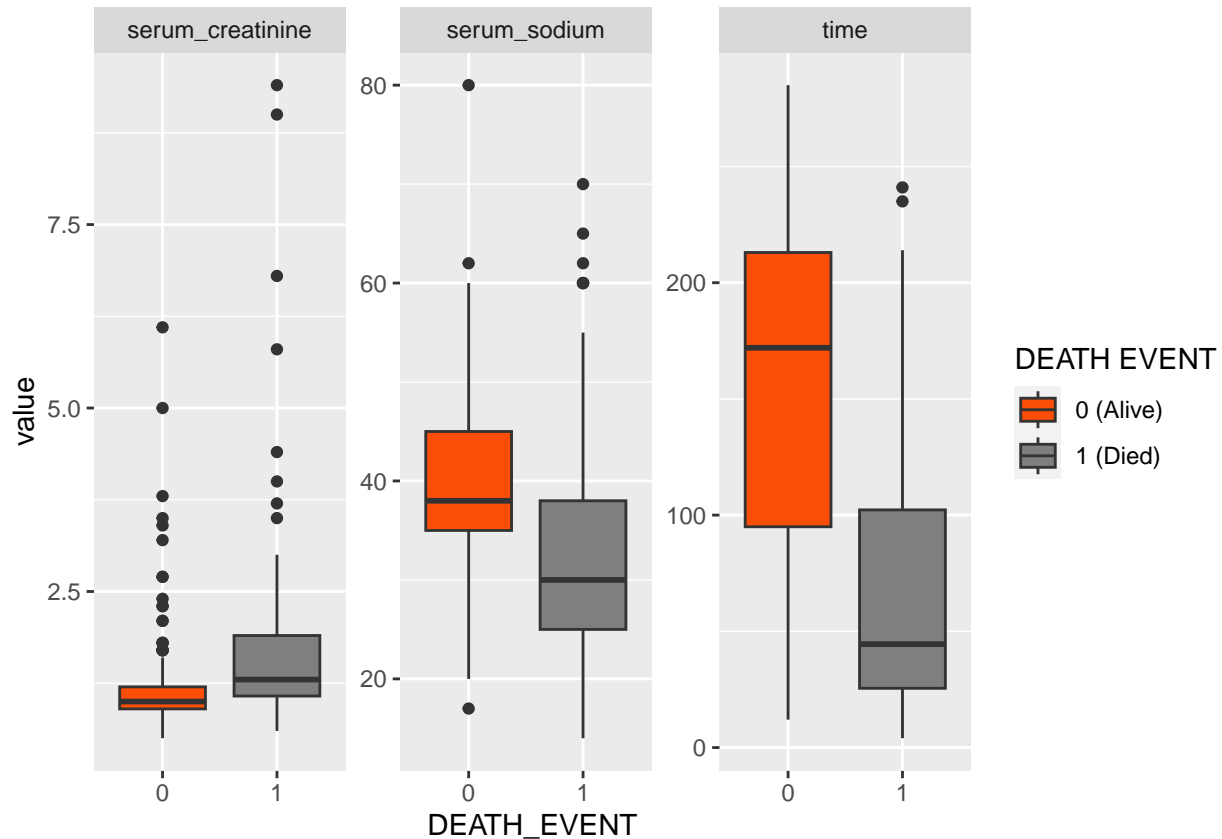
Distribution analysis

RELATIONSHIP OF PREDICTOR VARIABLES WITH DEATH EVENT COLUMN

```
# Box plot of age, creatinine phosphokinase, ejection fraction, platelets and death event
palette_ro <- c("yellow", "#FC4E07")
df1 <- melt(Heart_disease[,c(1,3,5,7,13)], id.var = "DEATH_EVENT")
ggplot(data = df1, aes(x=DEATH_EVENT, y=value)) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH_EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  geom_boxplot(aes(fill=DEATH_EVENT)) + facet_wrap(~variable, scales="free")
```



```
# Box plot of serum creatinine, serum sodium, time and death event
df2 <- melt(Heart_disease[,c(8,9,12,13)], id.var = "DEATH_EVENT")
ggplot(data = df2, aes(x=DEATH_EVENT, y=value)) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH_EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  geom_boxplot(aes(fill=DEATH_EVENT)) + facet_wrap(~variable, scales="free")
```

```
p11 = ggplot(data = Heart_disease, aes(x=anaemia)) +
  geom_bar(position = 'dodge', aes(fill=DEATH_EVENT), color = 'black') +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  labs(
    title = 'RELATIONSHIP OF ANAEMIA WITH DEATH EVENT'
  ) +
  scale_x_discrete(labels = c("No anaemia", "Anaemia"))

p22 = ggplot(data = Heart_disease, aes(x=diabetes)) +
  geom_bar(position = 'dodge', aes(fill=DEATH_EVENT), color = 'black') +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  labs(
    title = 'RELATIONSHIP OF DIABETES WITH DEATH EVENT'
  ) +
  scale_x_discrete(labels = c("No diabetes", "Diabetes"))

p33 = ggplot(data = Heart_disease, aes(x=high_blood_pressure)) +
  geom_bar(position = 'dodge', aes(fill=DEATH_EVENT), color = 'black') +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  labs(
```

```

    title = 'RELATIONSHIP OF DIABETES WITH DEATH EVENT'
  ) +
  scale_x_discrete(labels = c("No high BP", "High BP"))

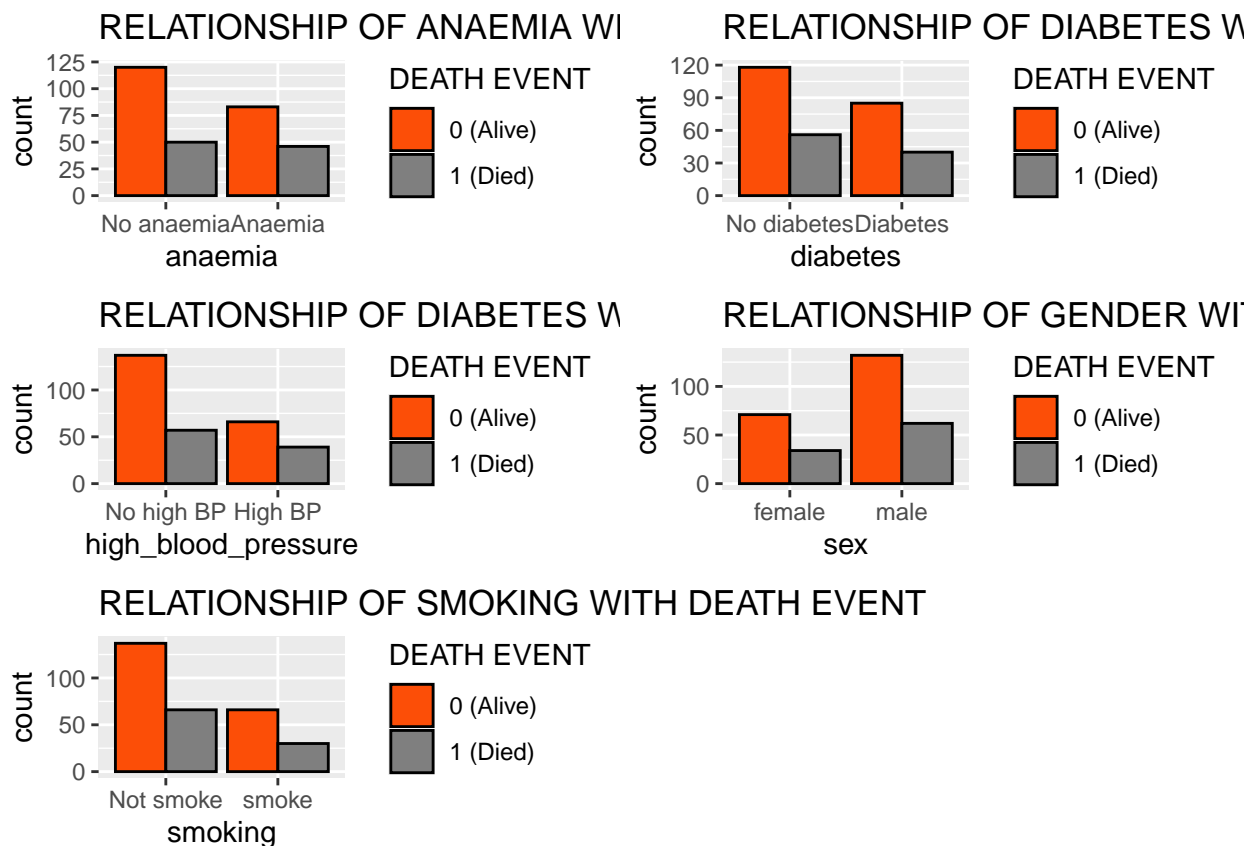
p44 = ggplot(data = Heart_disease, aes(x=sex)) +
  geom_bar(position = 'dodge', aes(fill=DEATH_EVENT), color = 'black') +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +

  labs(
    title = 'RELATIONSHIP OF GENDER WITH DEATH EVENT'
  ) +
  scale_x_discrete(labels = c("female", "male"))

p55 = ggplot(data = Heart_disease, aes(x=smoking)) +
  geom_bar(position = 'dodge', aes(fill=DEATH_EVENT), color = 'black') +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +

  labs(
    title = 'RELATIONSHIP OF SMOKING WITH DEATH EVENT'
  ) +
  scale_x_discrete(labels = c("Not smoke", "smoke"))
grid.arrange(p11, p22, p33, p44, p55, nrow = 3, ncol = 2)

```



We can observe many As we see in the boxplots creatinine phosphokinase and serum creatinine has high



variance and positively skewed. So we should carefully observe them in the histogram and density plots and also calculate the skewness and if necessary scale them.

Age distribution

```
#Histogram - age
palette_ro <- c("yellow", "#FC4E07")
p1 = ggplot(Heart_disease,aes(x=age)) +
  geom_histogram(aes(y = ..density..), binwidth = 5, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.25) +
  scale_x_continuous(breaks = seq(40, 100, 10)) +
  geom_vline(xintercept = median(Heart_disease$age), linetype="longdash", colour = "blue") +
  labs(x="Age of patient",
       y="Density",
       title="age distribution")

p2 = ggplot(Heart_disease, aes(x = age, fill = DEATH_EVENT)) +
  geom_density(aes(age,fill=DEATH_EVENT),alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  scale_x_continuous(breaks = seq(40, 100, 10)) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$age),
            linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$age),
            linetype="longdash", colour = "orange") +
  labs(x="Age of patient",
       y="Density",
       title="Relationship between age and DEATH EVENT")

grid.arrange(p1, p2, nrow=2)
```



The insights we get from above histogram and density for age distribution are: * The age of the patients was highest around 60 years and you can observe that, younger the age, the density plot of survival is high and as the age increases density plot of death is more. After the age 70 the density plot is reversed.

creatinine phosphokinase distribution

```
#Histogram - creatinine phosphokinase
palette_ro <- c("yellow", "#FC4E07")
p3 = ggplot(Heart_disease,aes(x=creatinine_phosphokinase)) +
  geom_histogram(aes(y = ..density..), binwidth = 100, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  geom_vline(xintercept = median(Heart_disease$creatinine_phosphokinase),
    linetype="longdash", colour = "blue") +
  labs(x="Level of the CPK enzyme in the blood (mcg/L)",
    y="Density",
    title="creatinine phosphokinase distribution")

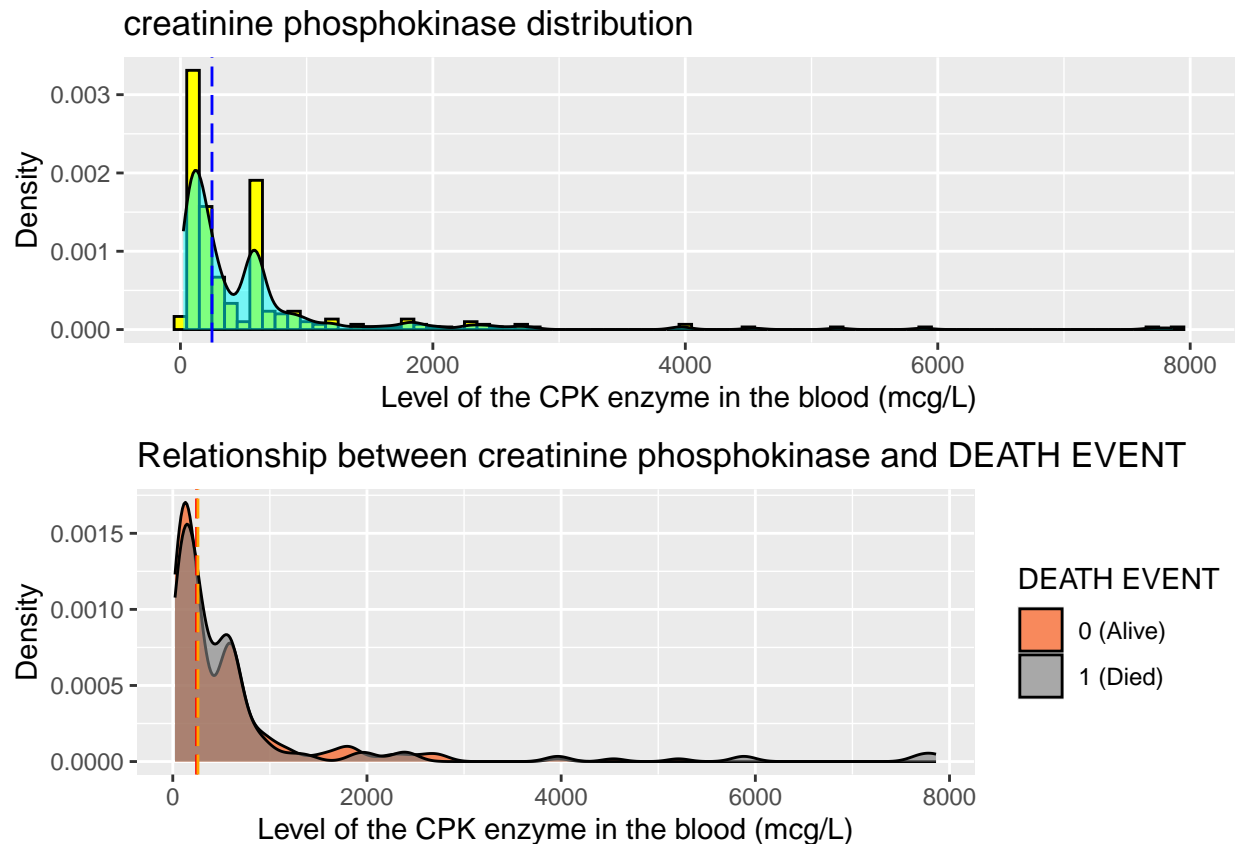
p4 = ggplot(Heart_disease, aes(x = creatinine_phosphokinase, fill = DEATH_EVENT)) +
  geom_density(aes(creatinine_phosphokinase,fill=DEATH_EVENT),alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
    name = "DEATH EVENT",
    labels = c("0 (Alive)", "1 (Died)")) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$creatinine_phosphokinase),
    linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$creatinine_phosphokinase),
```

```

    linetype="longdash", colour = "orange") +
  labs(x="Level of the CPK enzyme in the blood (mcg/L)",
       y="Density",
       title="Relationship between creatinine phosphokinase and DEATH EVENT")

grid.arrange(p3, p4, nrow=2)

```



The insights we get from above histogram and density for creatinine phosphokinase distribution are: * The distribution is heavily skewed to one side and we need to calculate the skewness and scale it accordingly.

ejection fraction distribution


```

#Histogram - ejection fraction
palette_ro <- c("yellow", "#FC4E07")
p5 = ggplot(Heart_disease,aes(x = ejection_fraction)) +
  geom_histogram(aes(y = ..density..), binwidth = 1, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  geom_vline(xintercept = median(Heart_disease$ejection_fraction), linetype="longdash",
             colour = "blue") +
  scale_x_continuous(breaks = seq(10, 80, 10)) +
  labs(x="Percentage of blood leaving the heart at each contraction",
       y="Density",
       title="ejection fraction distribution")

```

```
p6 = ggplot(Heart_disease, aes(x = ejection_fraction, fill = DEATH_EVENT)) +
  geom_density(aes(ejection_fraction, fill=DEATH_EVENT), alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$ejection_fraction),
             linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$ejection_fraction),
             linetype="longdash", colour = "orange") +
  scale_x_continuous(breaks = seq(10, 80, 10)) +
  labs(x="Percentage of blood leaving the heart at each contraction",
       y="Density",
       title="Relationship between ejection fraction and DEATH EVENT")

grid.arrange(p5, p6, nrow=2)
```



The insights we get from above histogram and density for creatinine phosphokinase distribution are: * The distribution is discrete. The ejection fraction for an normal person will be around 50-60 and if it is higher also, there will be no problem. But, if the ejection fraction falls to 40 and below there is a higher chances of heart failure. The same trend we can observe from the graph.

platelets distribution

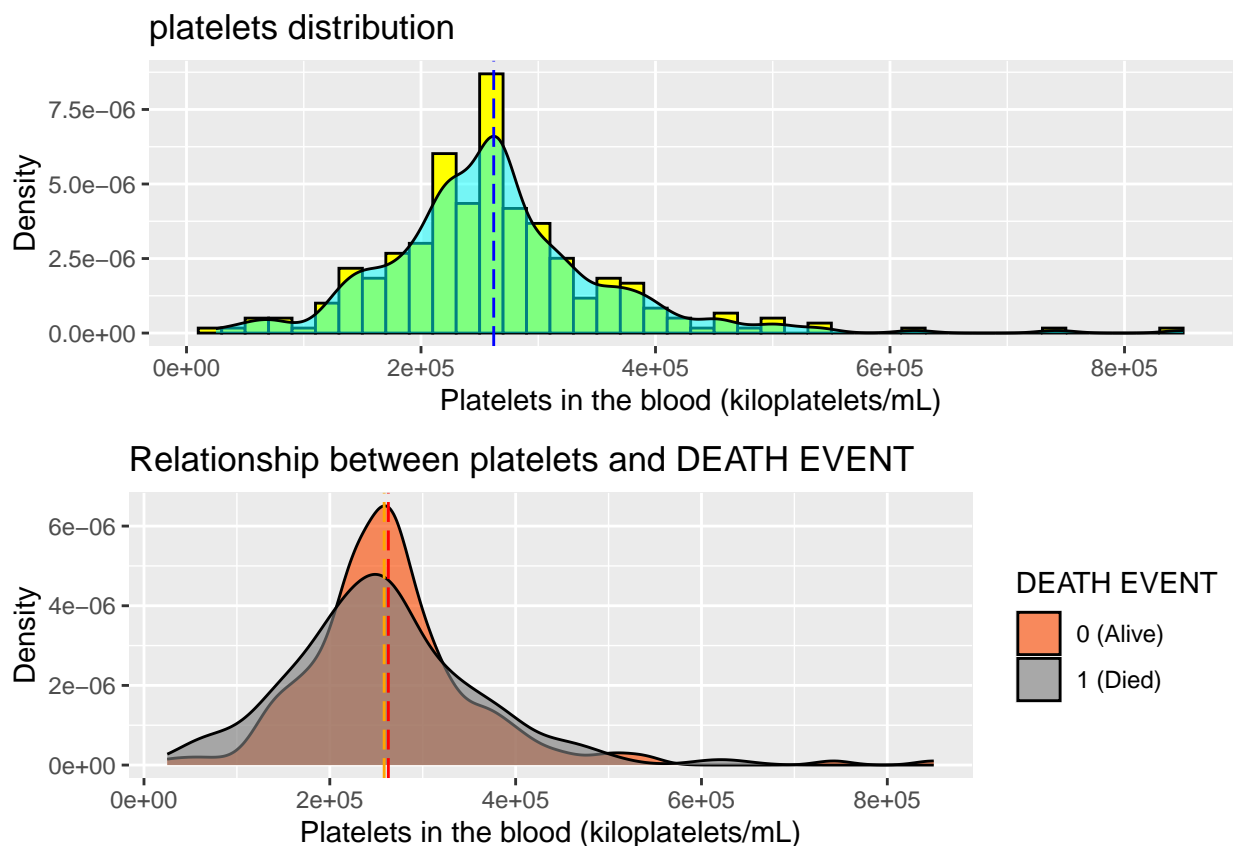
```

#Histogram - platelets
palette_ro <- c("yellow", "#FC4E07")
p7 = ggplot(Heart_disease,aes(x = platelets)) +
  geom_histogram(aes(y = ..density..), binwidth = 20000, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  geom_vline(xintercept = median(Heart_disease$platelets), linetype="longdash", colour = "blue") +
  labs(x="Platelets in the blood (kiloplatelets/mL)",
       y="Density",
       title="platelets distribution")

p8 = ggplot(Heart_disease, aes(x = platelets, fill = DEATH_EVENT)) +
  geom_density(aes(platelets,fill=DEATH_EVENT),alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$platelets),
             linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$platelets),
             linetype="longdash", colour = "orange") +
  labs(x="Platelets in the blood (kiloplatelets/mL)",
       y="Density",
       title="Relationship between platelets and DEATH EVENT")

grid.arrange(p7, p8, nrow=2)

```



The insights we get from above histogram and density for platelets distribution are: * The distribution is

symmetric , survivals have the highest platelets.


serum creatinine distribution

```
#Histogram - serum creatinine
palette_ro <- c("yellow", "#FC4E07")
p9 = ggplot(Heart_disease,aes(x = serum_creatinine )) +
  geom_histogram(aes(y = ..density..), binwidth = 0.2, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  geom_vline(xintercept = median(Heart_disease$serum_creatinine), linetype="longdash",
    colour = "blue") +
  labs(x="Level of serum creatinine in the blood (mg/dL)",
    y="Density",
    title="serum creatinine distribution")


p10 = ggplot(Heart_disease, aes(x = serum_creatinine, fill = DEATH_EVENT)) +
  geom_density(aes(serum_creatinine,fill=DEATH_EVENT),alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
    name = "DEATH EVENT",
    labels = c("0 (Alive)", "1 (Died)")) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$serum_creatinine),
    linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$serum_creatinine),
    linetype="longdash", colour = "orange") +
  labs(x="Level of serum creatinine in the blood (mg/dL)",
    y="Density",
    title="Relationship between serum creatinine and DEATH EVENT")

grid.arrange(p9, p10, nrow=2)
```


serum creatinine distribution



Relationship between serum creatinine and DEATH EVENT



The insights we get from above histogram and density for serum creatinine distribution are: * The distribution is highly skewed and will be scaled later in this research, for survivals the value is around the median and the patients who have 1.5mg/dL have high risk of heart failure.

serum sodium distribution

```
#Histogram - serum sodium
palette_ro <- c("yellow", "#FC4E07")
p11 = ggplot(Heart_disease,aes(x = serum_sodium)) +
  geom_histogram(aes(y = ..density..), binwidth = 1, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  scale_x_continuous(breaks = seq(100, 150, 10)) +
  geom_vline(xintercept = median(Heart_disease$serum_sodium), linetype="longdash",
    colour = "blue") +
  labs(x="Level of serum sodium in the blood (mEq/L)",
    y="Density",
    title="serum sodium distribution")

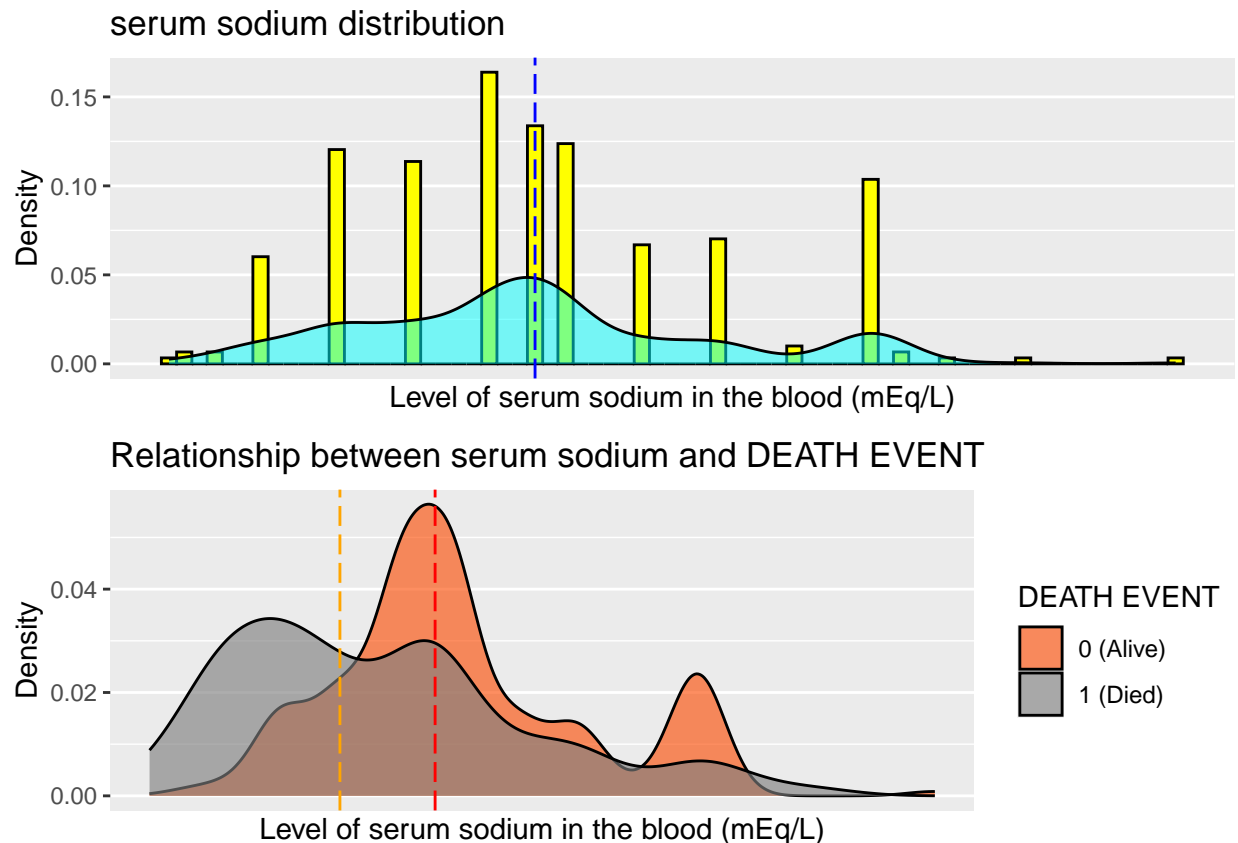
p12 = ggplot(Heart_disease, aes(x = serum_sodium, fill = DEATH_EVENT)) +
  geom_density(aes(serum_sodium,fill=DEATH_EVENT),alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
    name = "DEATH EVENT",
    labels = c("0 (Alive)", "1 (Died)")) +
  scale_x_continuous(breaks = seq(100, 150, 10)) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$serum_sodium),
```

```

    linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$serum_sodium),
    linetype="longdash", colour = "orange") +
  labs(x="Level of serum sodium in the blood (mEq/L)",
    y="Density",
    title="Relationship between serum sodium and DEATH EVENT")

grid.arrange(p11, p12, nrow=2)

```



The insights we get from above histogram and density for serum sodium distribution are: * The survival of patients is more around the median and the value of deaths get's lower when the level of serum sodium increases.

Time distribution

```

#Histogram - Time
palette_ro <- c("yellow", "#FC4E07")
p13 = ggplot(Heart_disease,aes(x = time)) +
  geom_histogram(aes(y = ..density..), binwidth = 10, fill = palette_ro[1],color="black") +
  geom_density(adjust=.8, fill="cyan",color="black", alpha=0.5) +
  scale_x_continuous(breaks = seq(0, 300, 50)) +
  geom_vline(xintercept = median(Heart_disease$time), linetype="longdash",
    colour = "blue") +
  labs(x="Follow-up period (days)",

```

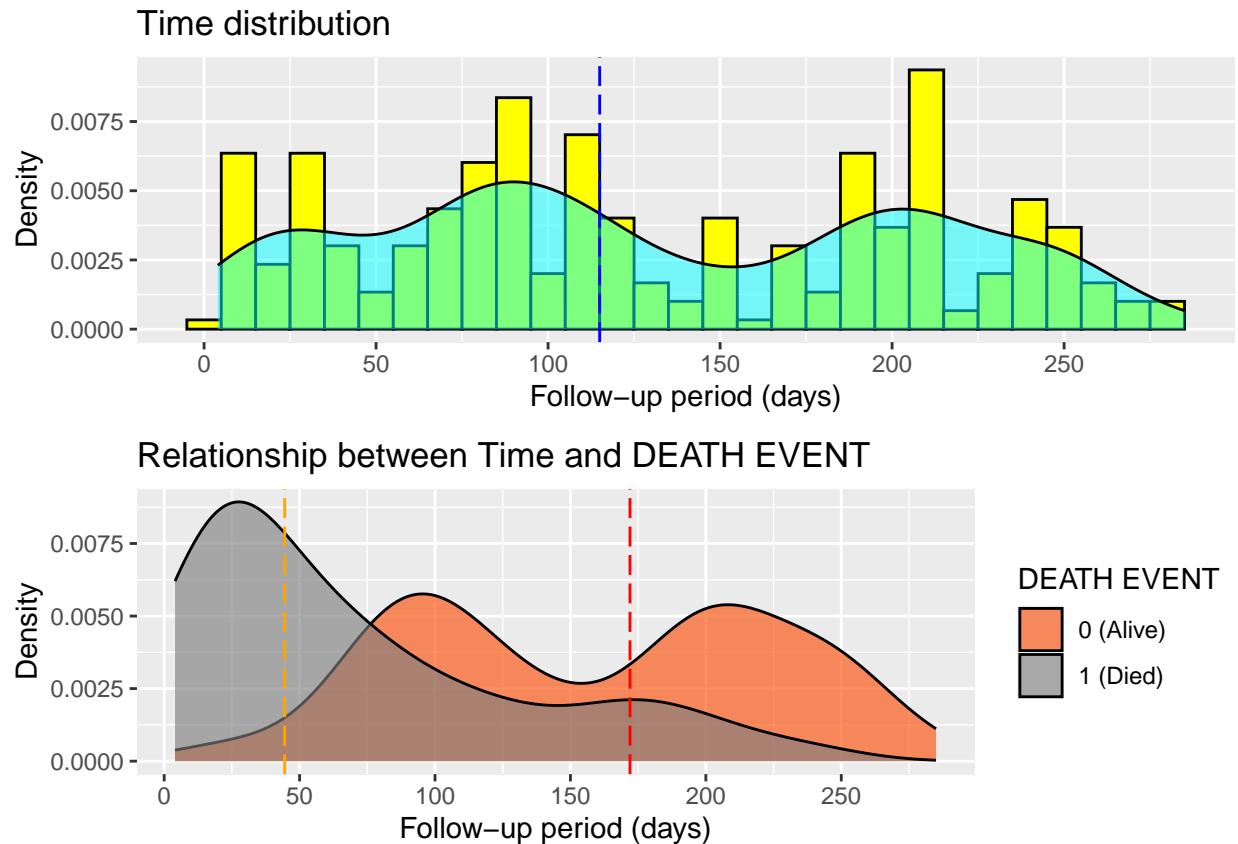
```

y="Density",
title="Time distribution")

p14 = ggplot(Heart_disease, aes(x = time, fill = DEATH_EVENT)) +
  geom_density(aes(time, fill=DEATH_EVENT), alpha=0.64) +
  scale_fill_manual(values = c(palette_ro[2], palette_ro[7]),
                    name = "DEATH EVENT",
                    labels = c("0 (Alive)", "1 (Died)")) +
  scale_x_continuous(breaks = seq(0, 300, 50)) +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 0)$time),
             linetype="longdash", colour = "red") +
  geom_vline(xintercept = median(filter(Heart_disease, DEATH_EVENT == 1)$time),
             linetype="longdash", colour = "orange") +
  labs(x="Follow-up period (days)",
       y="Density",
       title="Relationship between Time and DEATH EVENT")

grid.arrange(p13, p14, nrow=2)

```



The insights we get from above histogram and density for time distribution are: * The patients who have more follow-up days are likely to have higher chances of survival and the patients who have less than 60 follow-up days have higher chances of death. So, the more you follow-up on you regular health check-up's the more you have probability to survive.

From the above density and histogram plots, we can observe the data is moderately skewed, it may violate the assumptions of the model, leading to biased results or inaccurate predictions. Now let's calculate the

skewness of the dataset and observe which has zero, negative and positive skewness and transform them.

skewness of dataset (All numeric variables)

```
#using sapply function to find skewness of all numeric variables
numeric_cols <- Heart_disease %>% select_if(is.numeric)
skewness_values <- sapply(numeric_cols, skewness)
skewness_df <- data.frame(variable = names(skewness_values), skewness = skewness_values)
skewness_df$index <- row.names(skewness_df)
rownames(skewness_df) <- NULL # Remove row names
skewness_df <- skewness_df[, c("variable", "skewness")] # Rearrange columns
skewness_df
```

```
##           variable  skewness
## 1              age 0.4203739
## 2 creatinine_phosphokinase 4.4406886
## 3      ejection_fraction 0.5525927
## 4             platelets 1.4549746
## 5      serum_creatinine 4.4336102
## 6      serum_sodium 0.5525927
## 7              time 0.1271606
```

From the above dataframe we can observe that the variables creatinine_phosphokinase, serum_creatinine and possibly platelets have a high degree of skewness, which could potentially affect the performance of some statistical models. The results might improve if the values are scaled. Let's compare both scaled and un-scaled data below by fitting into different models.

Scaling the data (Data Normalization)

```
scaled_heart_disease <- Heart_disease
scaled_heart_disease[, c("age", "creatinine_phosphokinase", "ejection_fraction",
                        "platelets", "serum_creatinine",
                        "serum_sodium", "time")] <-
  scale(scaled_heart_disease[, c("age", "creatinine_phosphokinase", "ejection_fraction",
                                "platelets", "serum_creatinine",
                                "serum_sodium", "time")])

summary(scaled_heart_disease)
```

```
##      age      anaemia creatinine_phosphokinase diabetes
## Min.   :-1.75170    0:170   Min.   :-0.575952      0:174
## 1st Qu.: -0.82693    1:129   1st Qu.: -0.479589      1:125
## Median :-0.07029              Median :-0.342001
## Mean   : 0.00000              Mean   : 0.000000
## 3rd Qu.: 0.77041              3rd Qu.: 0.000165
## Max.    : 2.87217              Max.    : 7.502063
## ejection_fraction high_blood_pressure platelets      serum_creatinine
## Min.   :-2.034976    0:194   Min.   :-2.43607   Min.   :-0.864061
## 1st Qu.: -0.683035    1:105   1st Qu.: -0.52000   1st Qu.: -0.477404
```

```
## Median :-0.007065          Median :-0.01388   Median :-0.284076
## Mean  : 0.000000          Mean  : 0.00000   Mean  : 0.000000
## 3rd Qu.: 0.584409          3rd Qu.: 0.41043   3rd Qu.: 0.005916
## Max.   : 3.541779          Max.   : 5.99812   Max.   : 7.739045
## serum_sodium      sex      smoking      time      DEATH_EVENT
## Min.   :-2.034976    0:105    0:203    Min.   :-1.6268    0:203
## 1st Qu.: -0.683035    1:194    1: 96    1st Qu.: -0.7378    1: 96
## Median :-0.007065          Median :-0.1966
## Mean   : 0.000000          Mean   : 0.0000
## 3rd Qu.: 0.584409          3rd Qu.: 0.9372
## Max.   : 3.541779          Max.   : 1.9937
```

```
Heart_disease1 = subset(Heart_disease, select = -c(serum_sodium))

scaled_heart_disease11 <- Heart_disease1
scaled_heart_disease11[, c("age", "creatinine_phosphokinase", "ejection_fraction",
                           "platelets", "serum_creatinine",
                           "time")] <-
  scale(scaled_heart_disease11[, c("age", "creatinine_phosphokinase", "ejection_fraction",
                                   "platelets", "serum_creatinine",
                                   "time")])
```

Now you can observe that the scaled data has a mean of zero and a standard deviation of one.

Splitting the dataset into training and testing into 70% and 30%

```
set.seed(1)
train <- sample(nrow(Heart_disease), 0.7 * nrow(Heart_disease))
training_set_unscaled <- Heart_disease[train, ]
testing_set_unscaled <- Heart_disease[-train, ]

train_2 <- sample(nrow(scaled_heart_disease), 0.7 * nrow(scaled_heart_disease))
training_set_scaled <- scaled_heart_disease[train_2, ]
testing_set_scaled <- scaled_heart_disease[-train_2, ]

Heart_disease1 = subset(Heart_disease, select = -c(serum_sodium))

set.seed(1)
train11 <- sample(nrow(Heart_disease1), 0.7 * nrow(Heart_disease1))
training_set_unscaled11 <- Heart_disease1[train11, ]
testing_set_unscaled11 <- Heart_disease1[-train11, ]

train_22 <- sample(nrow(scaled_heart_disease11), 0.7 * nrow(scaled_heart_disease11))
training_set_scaled11 <- scaled_heart_disease11[train_22, ]
testing_set_scaled11 <- scaled_heart_disease11[-train_22, ]
```

Now let's fit the dataset to different classification models with scaled and unscaled data and compare the results of each.

MODEL 1 - LOGISTIC REGRESSION

MODEL 1.a -> LOGISTIC REGRESSION - ALL VARIABLES - UNSCALED

```
#fitting the un-scaled data on logistic regression model
glm.all_unscaled = glm(DEATH_EVENT~.,data=training_set_unscaled,family="binomial")

#predicting on the testing un-scaled data
predict_test_all_unscaled <- factor(ifelse(predict(glm.all_unscaled, testing_set_unscaled ,
                                                    type ="response") > 0.5, "1", "0"))
summary(glm.all_unscaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ ., family = "binomial", data = training_set_unscaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0887  -0.6121  -0.2442   0.4879   2.6104
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    5.038e-01  1.528e+00   0.330  0.741634
## age            5.829e-02  1.724e-02   3.381  0.000722 ***
## anaemia1       -3.303e-01  4.289e-01  -0.770  0.441311
## creatinine_phosphokinase 1.326e-04  1.713e-04   0.774  0.438792
## diabetes1       2.549e-01  4.066e-01   0.627  0.530734
## ejection_fraction -7.143e-02  1.924e-02  -3.713  0.000205 ***
## high_blood_pressure1 -4.397e-02  4.156e-01  -0.106  0.915729
## platelets       -2.195e-06  2.069e-06  -1.061  0.288643
## serum_creatinine  5.823e-01  1.872e-01   3.110  0.001870 **
## serum_sodium      NA          NA      NA      NA
## sex1            -4.960e-01  4.679e-01  -1.060  0.289138
## smoking1        -3.080e-02  4.890e-01  -0.063  0.949783
## time            -1.925e-02  3.447e-03  -5.584  2.35e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 162.08  on 197  degrees of freedom
## AIC: 186.08
##
## Number of Fisher Scoring iterations: 5
```

The logistic regression has fitted to whole dataset with death event as a response variable. We got an accuracy of 88.9% when logistic regression is done on the un-scaled data and a true positive rate (recall) of 93.44%. From Fig. 15 it can be noted the coefficients (age, ejection fraction, serum creatinine, time) have asterisks next to their p-values, indicating that they are statistically significant at certain significance levels ($p < 0.05$), while others are not. The coefficients with the variables which are insignificant have the p value greater than 0.05 meaning they have weak evidence against the null hypothesis, and we do not reject null

hypothesis. The lower p-values are considered more statistically significant and are typically interpreted as having a stronger association with the response variable (Death Event). The others with no asterisk are insignificant variables and we can remove them for further models.

We can also see serum sodium has NA values in the estimate column, standard error column, and p-value column. Which means the variable is not included in the model and the variable was likely dropped during the selection process. We observed that it is perfect positive with ejection fraction predictor. The main reason for the model to not select serum sodium is, this variable is highly correlated with the other predictor which is already included in the model.

Confusion matrix for full unscaled model

```
confusionMatrix(predict_test_all_unscaled,testing_set_unscaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 57   6
##           1   4 23
##
##           Accuracy : 0.8889
##           95% CI : (0.8051, 0.9454)
##       No Information Rate : 0.6778
##       P-Value [Acc > NIR] : 2.826e-06
##
##           Kappa : 0.7409
##
##  Mcnemar's Test P-Value : 0.7518
##
##           Sensitivity : 0.9344
##           Specificity : 0.7931
##           Pos Pred Value : 0.9048
##           Neg Pred Value : 0.8519
##           Prevalence : 0.6778
##           Detection Rate : 0.6333
##       Detection Prevalence : 0.7000
##           Balanced Accuracy : 0.8638
##
##           'Positive' Class : 0
##
```

We got an accuracy of 88.9% when logistic regression is done on the un-scaled data.

The confusion matrix obtained for un-scaled data is as follows: True Positives : 57 False Positives: 6 True Negatives : 23 False Negatives: 4

The balanced accuracy got by the logistic regression on the given dataset is 86.38%. It is the mean of sensitivity and Specificity.

Balanced Accuracy = Sensitivity + Specificity / 2 => 0.9344 + 0.7931 / 2 => 0.8638

Specificity value is 79.31% (True Negative Rate) Sensitivity value is 93.44% (True Positive Rate)

MODEL 1.b -> LOGISTIC REGRESSION - ALL VARIABLES - SCALED

```
#fitting the scaled data on logistic regression model
glm.all_scaled = glm(DEATH_EVENT~.,data=training_set_scaled,family="binomial")

#predicting on the testing scaled data
predict_test_all_scaled <- factor(ifelse(predict(glm.all_scaled, testing_set_scaled , type =
                                             "response") > 0.5, "1", "0"))
summary(glm.all_scaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ ., family = "binomial", data = training_set_scaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2618  -0.4682  -0.2058   0.3803   2.4580
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.02355    0.54624  -1.874 0.060954 .
## age             0.66366    0.24400   2.720 0.006530 **
## anaemia1       -0.35631    0.45968  -0.775 0.438272
## creatinine_phosphokinase 0.36738    0.23927   1.535 0.124689
## diabetes1       0.37462    0.43834   0.855 0.392750
## ejection_fraction -0.89478    0.24387  -3.669 0.000243 ***
## high_blood_pressure1 -0.01415    0.44318  -0.032 0.974523
## platelets       -0.13437    0.24610  -0.546 0.585079
## serum_creatinine  0.87110    0.21995   3.961 7.48e-05 ***
## serum_sodium      NA          NA      NA      NA
## sex1            -0.67105    0.51033  -1.315 0.188531
## smoking1         0.11796    0.51915   0.227 0.820253
## time            -1.84400    0.30653  -6.016 1.79e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 142.07  on 197  degrees of freedom
## AIC: 166.07
##
## Number of Fisher Scoring iterations: 6
```

Confusion matrix for full scaled model

```
confusionMatrix(predict_test_all_scaled,testing_set_scaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
```



```

##           Reference
## Prediction  0   1
##           0 56 10
##           1  5 19
##
##           Accuracy : 0.8333
##           95% CI : (0.74, 0.9036)
##           No Information Rate : 0.6778
##           P-Value [Acc > NIR] : 0.0006812
##
##           Kappa : 0.6004
##
## Mcnemar's Test P-Value : 0.3016996
##
##           Sensitivity : 0.9180
##           Specificity : 0.6552
##           Pos Pred Value : 0.8485
##           Neg Pred Value : 0.7917
##           Prevalence : 0.6778
##           Detection Rate : 0.6222
##           Detection Prevalence : 0.7333
##           Balanced Accuracy : 0.7866
##
##           'Positive' Class : 0
##

```

We got an accuracy of 83.3% when logistic regression is done on the scaled data.

The confusion matrix obtained for scaled data is as follows: True Positives : 56 False Positives: 10 True Negatives : 19 False Negatives: 5

The balanced accuracy got by the logistic regression on the given dataset is 78.66%. It is the mean of sensitivity and Specificity.

Balanced Accuracy = Sensitivity + Specificity / 2 => 0.9180 + 0.6552 / 2 => 0.7866

Specificity value is 65.52% (True Negative Rate) Sensitivity value is 91.80% (True Positive Rate)

As we can see that the un-scaled data when fitted for the logistic model we got an accuracy of 88.9% and when fitted with the scaled data we got an accuracy of 83.3%. It is possible that scaling the data may have reduced the accuracy of the logistic regression model compared to the un-scaled data. This could be due to a few reasons:

- Outliers: Scaling the data can sometimes amplify the effect of outliers, which can negatively impact the accuracy of the model.
- Non-linear relationships: If there are non-linear relationships between the predictors and the response, scaling the data may reduce the ability of the model to capture these relationships.

In general, It is good to experiment both scaled and un-scaled data and choose the method that has good accuracy. Here in our case we choose un-scaled data fitted for logistic model.

Since, We had perfect positive correlation between serum sodium, ejection fraction. Lets try removing one variable and check the accuracy. Now let's remove the serum sodium and observe whether it effects the model in any way for both scaled and unscaled data.

MODEL 1.c -> LOGISTIC REGRESSION - REMOVING SERUM SODIUM VARIABLE - UN-SCALED

```
#fitting the un-scaled data on logistic regression model (removing serum sodium)
glm.removing_serum_sodium_unscaled = glm(DEATH_EVENT~.-serum_sodium,data=training_set_unscaled,
                                          family="binomial")

#predicting on the testing un-scaled data (removing serum sodium)
predict_test_unscaled <- factor(ifelse(predict(glm.removing_serum_sodium_unscaled, testing_set_unscaled
                                              type ="response") > 0.5, "1", "0"))

summary(glm.removing_serum_sodium_unscaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ . - serum_sodium, family = "binomial",
##      data = training_set_unscaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0887  -0.6121  -0.2442   0.4879   2.6104
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.038e-01  1.528e+00   0.330  0.741634
## age              5.829e-02  1.724e-02   3.381  0.000722 ***
## anaemia1        -3.303e-01  4.289e-01  -0.770  0.441311
## creatinine_phosphokinase 1.326e-04  1.713e-04   0.774  0.438792
## diabetes1        2.549e-01  4.066e-01   0.627  0.530734
## ejection_fraction -7.143e-02  1.924e-02  -3.713  0.000205 ***
## high_blood_pressure1 -4.397e-02  4.156e-01  -0.106  0.915729
## platelets        -2.195e-06  2.069e-06  -1.061  0.288643
## serum_creatinine   5.823e-01  1.872e-01   3.110  0.001870 **
## sex1             -4.960e-01  4.679e-01  -1.060  0.289138
## smoking1         -3.080e-02  4.890e-01  -0.063  0.949783
## time            -1.925e-02  3.447e-03  -5.584  2.35e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 162.08  on 197  degrees of freedom
## AIC: 186.08
##
## Number of Fisher Scoring iterations: 5
```

Confusion matrix for removing serum sodium predictor - un-scaled

```
confusionMatrix(predict_test_unscaled,testing_set_unscaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction 0  1
##           0 57  6
##           1  4 23
##
##           Accuracy : 0.8889
##           95% CI : (0.8051, 0.9454)
##           No Information Rate : 0.6778
##           P-Value [Acc > NIR] : 2.826e-06
##
##           Kappa : 0.7409
##
## Mcnemar's Test P-Value : 0.7518
##
##           Sensitivity : 0.9344
##           Specificity : 0.7931
##           Pos Pred Value : 0.9048
##           Neg Pred Value : 0.8519
##           Prevalence : 0.6778
##           Detection Rate : 0.6333
##           Detection Prevalence : 0.7000
##           Balanced Accuracy : 0.8638
##
##           'Positive' Class : 0
##
```

MODEL 1.d -> LOGISTIC REGRESSION - REMOVING SERUM SODIUM VARIABLE - SCALED

```
#fitting the scaled data on logistic regression model (removing serum sodium)
glm.removing_serum_sodium_scaled = glm(DEATH_EVENT~.-serum_sodium,data=training_set_scaled,
                                       family="binomial")

#predicting on the testing scaled data (removing serum sodium)
predict_test_scaled <- factor(ifelse(predict(glm.removing_serum_sodium_scaled, testing_set_scaled ,
                                             type ="response") > 0.5, "1", "0"))
summary(glm.removing_serum_sodium_scaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ . - serum_sodium, family = "binomial",
##      data = training_set_scaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2618  -0.4682  -0.2058   0.3803   2.4580
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.02355    0.54624  -1.874 0.060954 .
## age           0.66366    0.24400   2.720 0.006530 **
```

```
## anaemia1          -0.35631    0.45968  -0.775  0.438272
## creatinine_phosphokinase  0.36738    0.23927   1.535  0.124689
## diabetes1         0.37462    0.43834   0.855  0.392750
## ejection_fraction  -0.89478    0.24387  -3.669  0.000243 ***
## high_blood_pressure1 -0.01415    0.44318  -0.032  0.974523
## platelets         -0.13437    0.24610  -0.546  0.585079
## serum_creatinine    0.87110    0.21995   3.961  7.48e-05 ***
## sex1              -0.67105    0.51033  -1.315  0.188531
## smoking1           0.11796    0.51915   0.227  0.820253
## time              -1.84400    0.30653  -6.016  1.79e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 142.07  on 197  degrees of freedom
## AIC: 166.07
##
## Number of Fisher Scoring iterations: 6
```

Confusion matrix for removing serum sodium predictor - scaled

```
confusionMatrix(predict_test_scaled,testing_set_scaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 56 10
##           1  5 19
##
##           Accuracy : 0.8333
##           95% CI : (0.74, 0.9036)
##    No Information Rate : 0.6778
##    P-Value [Acc > NIR] : 0.0006812
##
##           Kappa : 0.6004
##
## Mcnemar's Test P-Value : 0.3016996
##
##           Sensitivity : 0.9180
##           Specificity : 0.6552
##           Pos Pred Value : 0.8485
##           Neg Pred Value : 0.7917
##           Prevalence : 0.6778
##           Detection Rate : 0.6222
##    Detection Prevalence : 0.7333
##           Balanced Accuracy : 0.7866
##
##           'Positive' Class : 0
##
```

We got the same confusion matrix for the full model and when serum sodium is removed from both scaled and un-scaled data. We observe no difference when the predictor serum sodium is removed from both scaled and un-scaled data. The accuracies remains the same i.e; 83.3% and 88.89% respectively.

Now let's find the significant predictors and fit the model and compare the accuracies for both scaled and unscaled data and also compare the model with the above models.

FORWARD STEPWISE SELECTION

```
Heart_disease2 = subset(Heart_disease, select = -c(serum_sodium))
regfit.bwd=regsubsets(DEATH_EVENT~.,data=Heart_disease2,nvmax=12, method ="forward")
summary(regfit.bwd)
```

```
## Subset selection object
## Call: regsubsets.formula(DEATH_EVENT ~ ., data = Heart_disease2, nvmax = 12,
##      method = "forward")
## 11 Variables (and intercept)
##
```

		Forced in	Forced out
## age		FALSE	FALSE
## anaemia1		FALSE	FALSE
## creatinine_phosphokinase		FALSE	FALSE
## diabetes1		FALSE	FALSE
## ejection_fraction		FALSE	FALSE
## high_blood_pressure1		FALSE	FALSE
## platelets		FALSE	FALSE
## serum_creatinine		FALSE	FALSE
## sex1		FALSE	FALSE
## smoking1		FALSE	FALSE
## time		FALSE	FALSE

```
## 1 subsets of each size up to 11
## Selection Algorithm: forward
##
```

		age	anaemia1	creatinine_phosphokinase	diabetes1	ejection_fraction
## 1 (1)	" " " "	" "	" "	" "	" "	" "
## 2 (1)	" " " "	" "	" "	" "	" "	"*
## 3 (1)	" " " "	" "	" "	" "	" "	"*
## 4 (1)	"*" " "	" "	" "	" "	" "	"*
## 5 (1)	"*" " "	"*	" "	" "	" "	"*
## 6 (1)	"*" " "	"*	" "	" "	" "	"*
## 7 (1)	"*" " "	"*	" "	"*	" "	"*
## 8 (1)	"*" " "	"*	" "	"*	" "	"*
## 9 (1)	"*" " "	"*	" "	"*	" "	"*
## 10 (1)	"*" "*" "	"*	" "	"*	" "	"*
## 11 (1)	"*" "*" "	"*	" "	"*	" "	"*

```
##
```

		high_blood_pressure1	platelets	serum_creatinine	sex1	smoking1	time
## 1 (1)	" "	" "	" "	" "	" "	" "	"*
## 2 (1)	" "	" "	" "	" "	" "	" "	"*
## 3 (1)	" "	" "	"*	" "	" "	" "	"*
## 4 (1)	" "	" "	"*	" "	" "	" "	"*
## 5 (1)	" "	" "	"*	" "	" "	" "	"*
## 6 (1)	" "	" "	"*	"*	" "	" "	"*
## 7 (1)	" "	" "	"*	"*	" "	" "	"*
## 8 (1)	" "	"*	"*	"*	" "	" "	"*

```
## 9 ( 1 ) "*" "*" "*" "*" " " "*"
## 10 ( 1 ) "*" "*" "*" "*" " " "*"
## 11 ( 1 ) "*" "*" "*" "*" "*" "*" "
```

BACKWARD STEPWISE SELECTION

```
Heart_disease2 = subset(Heart_disease, select = -c(serum_sodium))
regfit.bwd=regsubsets(DEATH_EVENT~.,data=Heart_disease2,nvmax=12, method ="backward")
summary(regfit.bwd)
```

```
## Subset selection object
## Call: regsubsets.formula(DEATH_EVENT ~ ., data = Heart_disease2, nvmax = 12,
##      method = "backward")
## 11 Variables (and intercept)
##               Forced in Forced out
## age                FALSE      FALSE
## anaemia1           FALSE      FALSE
## creatinine_phosphokinase  FALSE      FALSE
## diabetes1          FALSE      FALSE
## ejection_fraction  FALSE      FALSE
## high_blood_pressure1 FALSE      FALSE
## platelets          FALSE      FALSE
## serum_creatinine   FALSE      FALSE
## sex1               FALSE      FALSE
## smoking1           FALSE      FALSE
## time               FALSE      FALSE
## 1 subsets of each size up to 11
## Selection Algorithm: backward
##      age anaemia1 creatinine_phosphokinase diabetes1 ejection_fraction
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " "*"
## 3 ( 1 ) " " " " " " "*"
## 4 ( 1 ) "*" " " " " " "*"
## 5 ( 1 ) "*" " " "*" " " "*"
## 6 ( 1 ) "*" " " "*" " " "*"
## 7 ( 1 ) "*" " " "*" " " "*"
## 8 ( 1 ) "*" " " "*" " " "*"
## 9 ( 1 ) "*" " " "*" " " "*"
## 10 ( 1 ) "*" "*" "*" " " "*"
## 11 ( 1 ) "*" "*" "*" " " "*"
##      high_blood_pressure1 platelets serum_creatinine sex1 smoking1 time
## 1 ( 1 ) " " " " " " " " "*"
## 2 ( 1 ) " " " " " " " " "*"
## 3 ( 1 ) " " " " "*" " " " "*"
## 4 ( 1 ) " " " " "*" " " " "*"
## 5 ( 1 ) " " " " "*" " " " "*"
## 6 ( 1 ) " " " " "*" " " " "*"
## 7 ( 1 ) " " " " "*" " " " "*"
## 8 ( 1 ) " " "*" "*" " " " " "*"
## 9 ( 1 ) "*" "*" "*" " " " " "*"
## 10 ( 1 ) "*" "*" "*" " " " " "*"
## 11 ( 1 ) "*" "*" "*" " " " " "
```

We see that using backward stepwise selection, the best one-variable model contains only time, and the best two-variable model additionally includes ejection_fraction. After observing the summary of the full model with the p-values of the predictors and the variable selection technique (backward), we can conclude that time, ejection_fraction, serum_creatinine and age are taken as the significant variables.

MODEL 1.e -> LOGISTIC REGRESSION - SIGNIFICANT VARIABLES - UN-SCALED

```
#fitting the un-scaled data on logistic regression model (only significant variables)
glm.significant_unscaled =
  glm(DEATH_EVENT~time+ejection_fraction+serum_creatinine+age,data=training_set_unscaled,
      family="binomial")

#predicting on the testing un-scaled data
predict_test_unscaled1 <- factor(ifelse(predict(glm.significant_unscaled, testing_set_unscaled,
                                             type = "response") > 0.5, "1", "0"))

summary(glm.significant_unscaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ time + ejection_fraction + serum_creatinine +
##      age, family = "binomial", data = training_set_unscaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9636  -0.6385  -0.2597   0.4775   2.7971
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.131529   1.183781  -0.111 0.911530
## time         -0.019215   0.003244  -5.923 3.16e-09 ***
## ejection_fraction -0.069250  0.018456  -3.752 0.000175 ***
## serum_creatinine  0.599730  0.180621   3.320 0.000899 ***
## age           0.052946   0.016381   3.232 0.001229 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 165.55  on 204  degrees of freedom
## AIC: 175.55
##
## Number of Fisher Scoring iterations: 5
```

Confusion matrix for significant predictors - un-scaled

```
confusionMatrix(predict_test_unscaled1,testing_set_unscaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction 0  1
##           0 57  6
##           1  4 23
##
##           Accuracy : 0.8889
##           95% CI : (0.8051, 0.9454)
##           No Information Rate : 0.6778
##           P-Value [Acc > NIR] : 2.826e-06
##
##           Kappa : 0.7409
##
## Mcnemar's Test P-Value : 0.7518
##
##           Sensitivity : 0.9344
##           Specificity : 0.7931
##           Pos Pred Value : 0.9048
##           Neg Pred Value : 0.8519
##           Prevalence : 0.6778
##           Detection Rate : 0.6333
##           Detection Prevalence : 0.7000
##           Balanced Accuracy : 0.8638
##
##           'Positive' Class : 0
##
```

MODEL 1.f -> LOGISTIC REGRESSION - SIGNIFICANT VARIABLES - SCALED

```
#fitting the scaled data on logistic regression model (only significant variables)
glm.significant_scaled =
  glm(DEATH_EVENT~age+ejection_fraction+serum_creatinine+time,data=training_set_scaled,
      family="binomial")

#predicting on the testing scaled data
predict_test_scaled1 <- factor(ifelse(predict(glm.significant_scaled, testing_set_scaled ,
      type ="response") > 0.5, "1", "0"))

summary(glm.significant_scaled)
```

```
##
## Call:
## glm(formula = DEATH_EVENT ~ age + ejection_fraction + serum_creatinine +
##      time, family = "binomial", data = training_set_scaled)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2894  -0.4975  -0.2151   0.4442   2.3156
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.3340     0.2465  -5.411 6.27e-08 ***
## age              0.5351     0.2249   2.380 0.017327 *
```



```
## ejection_fraction -0.8526      0.2293 -3.718 0.000201 ***
## serum_creatinine  0.8528      0.2120  4.023 5.74e-05 ***
## time              -1.7865      0.2854 -6.259 3.87e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 262.21  on 208  degrees of freedom
## Residual deviance: 148.29  on 204  degrees of freedom
## AIC: 158.29
##
## Number of Fisher Scoring iterations: 6
```

Confusion matrix for significant predictors -scaled

```
confusionMatrix(predict_test_scaled1,testing_set_scaled$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 54  9
##           1  7 20
##
##           Accuracy : 0.8222
##           95% CI : (0.7274, 0.8948)
##           No Information Rate : 0.6778
##           P-Value [Acc > NIR] : 0.001599
##
##           Kappa : 0.5855
##
## Mcnemar's Test P-Value : 0.802587
##
##           Sensitivity : 0.8852
##           Specificity : 0.6897
##           Pos Pred Value : 0.8571
##           Neg Pred Value : 0.7407
##           Prevalence : 0.6778
##           Detection Rate : 0.6000
##           Detection Prevalence : 0.7000
##           Balanced Accuracy : 0.7875
##
##           'Positive' Class : 0
##
```


We got the same confusion matrix and the calculations of the model's performance. So, the model with 4 significant variables (age, ejection_fraction, serum_creatinine and time) is selected as the best model with an accuracy of 88.89% and recall/ true positive rate/ sensitivity of 93.44%.

MODEL 2 - K-NEAREST NEIGHBORS

MODEL 2.a - K-NEAREST NEIGHBORS - ALL VARIABLES - UNSCALED

```
error<-rep(NA,6) # Placeholder
training_set_unscaled_knn = subset(training_set_unscaled, select =
                                   c(age,creatinine_phosphokinase,ejection_fraction,
                                     platelets,serum_creatinine,time))
testing_set_unscaled_knn = subset(testing_set_unscaled, select =
                                   c(age,creatinine_phosphokinase,ejection_fraction,platelets,
                                     serum_creatinine,time))

for(i in 1:6)
{
  knn.pred = knn(training_set_unscaled_knn,testing_set_unscaled_knn,
                  training_set_unscaled$DEATH_EVENT,k=i)
  error[i]=mean(knn.pred!=testing_set_unscaled$DEATH_EVENT)
}
plot(error,type="b",xlab="K",ylab="Test Error")
```



Finding best k-value

```
k = (loc=which.min(error))
k
```

```
## [1] 6
```

Confusion matrix for KNN - un-scaled (All variables)

```
knn.pred_unscaled=knn(train = training_set_unscaled_knn,test = testing_set_unscaled_knn,cl =
                      training_set_unscaled$DEATH_EVENT,k = k)
confusionMatrix(testing_set_unscaled$DEATH_EVENT, knn.pred_unscaled)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 48 13
##           1 25   4
##
##           Accuracy : 0.5778
##           95% CI : (0.4691, 0.6812)
##       No Information Rate : 0.8111
##       P-Value [Acc > NIR] : 1.00000
##
##           Kappa : -0.0843
##
##  McNemar's Test P-Value : 0.07435
##
##           Sensitivity : 0.6575
##           Specificity : 0.2353
##           Pos Pred Value : 0.7869
##           Neg Pred Value : 0.1379
##           Prevalence : 0.8111
##           Detection Rate : 0.5333
##       Detection Prevalence : 0.6778
##           Balanced Accuracy : 0.4464
##
##           'Positive' Class : 0
##
```

From the un-scaled model fitted to KNN, the obtained accuracy is 65.56% for full variable (continuous) model.


MODEL 2.b - K-NEAREST NEIGHBORS - ALL VARIABLES - SCALED

```
error<-rep(NA,6) # Placeholder
training_set_scaled_knn = subset(training_set_unscaled, select =
                                c(age,creatinine_phosphokinase,ejection_fraction,
                                platelets,serum_creatinine,time))
testing_set_scaled_knn = subset(testing_set_unscaled, select =
```

```

c(age,creatinine_phosphokinase,ejection_fraction,
  platelets,serum_creatinine,time))
for(i in 1:6)
{
  knn.pred=knn(training_set_scaled_knn,testing_set_scaled_knn,training_set_scaled$DEATH_EVENT,
               k=i)
  error[i]=mean(knn.pred!=testing_set_scaled$DEATH_EVENT)
}
plot(error,type="b",xlab="K",ylab="Test Error")

```



Finding best k-value

```

k = (loc=which.min(error))
k

```

```
## [1] 3
```

Confusion matrix for KNN - scaled (All variables)

```

knn.pred_scaled=knn(train = training_set_scaled_knn,test = testing_set_scaled_knn,
                    cl = training_set_scaled$DEATH_EVENT,k = k)
confusionMatrix(testing_set_scaled$DEATH_EVENT, knn.pred_scaled)

```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 51 10
##           1 24   5
##
##           Accuracy : 0.6222
##           95% CI : (0.5138, 0.7223)
##       No Information Rate : 0.8333
##       P-Value [Acc > NIR] : 1.00000
##
##           Kappa : 0.0097
##
##  McNemar's Test P-Value : 0.02578
##
##           Sensitivity : 0.6800
##           Specificity : 0.3333
##       Pos Pred Value : 0.8361
##       Neg Pred Value : 0.1724
##           Prevalence : 0.8333
##       Detection Rate : 0.5667
##       Detection Prevalence : 0.6778
##       Balanced Accuracy : 0.5067
##
##       'Positive' Class : 0
##
```

From the scaled model fitted to KNN, the obtained accuracy is 87.78% for full variable (continuous) model.

MODEL 2.c - K-NEAREST NEIGHBORS - SIGNIFICANT VARIABLES - UN-SCALED

```
error<-rep(NA,4) # Placeholder
training_set_unscaled_knn1 = subset(training_set_unscaled, select =
                                   c(age,ejection_fraction,
                                     serum_creatinine,time))
testing_set_unscaled_knn1 = subset(testing_set_unscaled, select =
                                   c(age,ejection_fraction,
                                     serum_creatinine,time))

for(i in 1:4)
{
  knn.pred = knn(training_set_unscaled_knn1,testing_set_unscaled_knn1,
                 training_set_unscaled$DEATH_EVENT,k=i)
  error[i]=mean(knn.pred!=testing_set_unscaled$DEATH_EVENT)
}
plot(error,type="b",xlab="K",ylab="Test Error")
```



Finding best k-value

```
k = (loc=which.min(error))
k
```

```
## [1] 1
```

Confusion matrix for KNN - un-scaled (significant variables)

```
knn.pred_unscaled1 = knn(train = training_set_unscaled_knn1, test = testing_set_unscaled_knn1,
                          cl = training_set_unscaled$DEATH_EVENT, k = k)
confusionMatrix(testing_set_unscaled$DEATH_EVENT, knn.pred_unscaled1)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0   1
```

```
##           0 57  4
```

```
##           1  9 20
```

```
##
```

```
##           Accuracy : 0.8556
```


```
##          95% CI : (0.7657, 0.9208)
##    No Information Rate : 0.7333
##    P-Value [Acc > NIR] : 0.004229
##
##          Kappa : 0.6536
##
##    McNemar's Test P-Value : 0.267257
##
##          Sensitivity : 0.8636
##          Specificity : 0.8333
##    Pos Pred Value : 0.9344
##    Neg Pred Value : 0.6897
##          Prevalence : 0.7333
##    Detection Rate : 0.6333
##    Detection Prevalence : 0.6778
##    Balanced Accuracy : 0.8485
##
##    'Positive' Class : 0
##
```

From the un-scaled model fitted to KNN, the obtained accuracy is 85.56% for significant variables.

MODEL 2.d - K-NEAREST NEIGHBORS - SIGNIFICANT VARIABLES - SCALED

```
error<-rep(NA,4) # Placeholder
training_set_scaled_knn1 = subset(training_set_scaled, select =
                                c(age,ejection_fraction,
                                  serum_creatinine,time))
testing_set_scaled_knn1 = subset(testing_set_scaled, select =
                                c(age,ejection_fraction,
                                  serum_creatinine,time))

for(i in 1:4)
{
  knn.pred = knn(training_set_scaled_knn1,testing_set_scaled_knn1,training_set_scaled$DEATH_EVENT,
                 k=i)
  error[i]=mean(knn.pred!=testing_set_scaled$DEATH_EVENT)
}
plot(error,type="b",xlab="K",ylab="Test Error")
```



Finding best k-value

```
k = (loc=which.min(error))
k
```

```
## [1] 3
```

Confusion matrix for KNN - scaled (significant variables)

```
knn.pred_scaled1 = knn(train = training_set_scaled_knn1, test = testing_set_scaled_knn1, cl = training_set_scaled_knn1$DEATH_EVENT)
confusionMatrix(testing_set_scaled$DEATH_EVENT, knn.pred_scaled1)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 52  9
##           1  8 21
##
##               Accuracy : 0.8111
##               95% CI : (0.7149, 0.8859)
```



```
##      No Information Rate : 0.6667
##      P-Value [Acc > NIR] : 0.001794
##
##              Kappa : 0.5714
##
##  McNemar's Test P-Value : 1.000000
##
##      Sensitivity : 0.8667
##      Specificity : 0.7000
##      Pos Pred Value : 0.8525
##      Neg Pred Value : 0.7241
##      Prevalence : 0.6667
##      Detection Rate : 0.5778
##      Detection Prevalence : 0.6778
##      Balanced Accuracy : 0.7833
##
##      'Positive' Class : 0
##
```

From the scaled model fitted to KNN, the obtained accuracy is 81.11% for significant variables.

We can observe the KNN algorithm performs well on un-scaled data compared to scaled data. So, the model with un-scaled data with significant variables is selected as best model in KNN with 85.56 accuracy and will compare it with other models later in this research.

MODEL 3 - DECISION TREES


MODEL 3.a - DECISION TREES - ALL VARIABLES - UN-SCALED

```
tree.Heart_disease_unscaled = tree(DEATH_EVENT~.,training_set_unscaled)
summary(tree.Heart_disease_unscaled)
```

```
##
## Classification tree:
## tree(formula = DEATH_EVENT ~ ., data = training_set_unscaled)
## Variables actually used in tree construction:
## [1] "time"                "platelets"
## [3] "creatinine_phosphokinase" "ejection_fraction"
## [5] "age"                 "serum_creatinine"
## Number of terminal nodes: 18
## Residual mean deviance: 0.3502 = 66.89 / 191
## Misclassification error rate: 0.08134 = 17 / 209
```

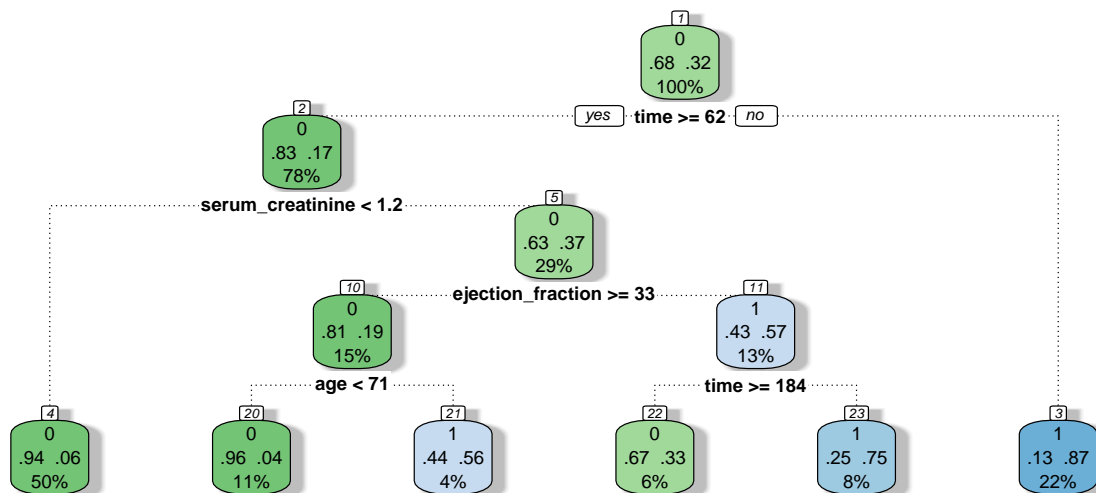
The summary() indicates that 6 of the variables are used in constructing the tree.

```
plot(tree.Heart_disease_unscaled)
text(tree.Heart_disease_unscaled,pretty=0)
```



Plotting the tree

```
#Better tree visualization using rpart package
tree.Heart_disease_unscaled = rpart(DEATH_EVENT~., training_set_unscaled)
fancyRpartPlot(tree.Heart_disease_unscaled)
```



Rattle 2023-May-16 13:31:19 vishaypaka

The type="class" argument specifies that the predicted values should be the class labels rather than probabilities.

We can come up with research question based on the tree: If a patient of age 78 admitted in hospital and stayed more than 70 days and the level of serum creatinine in his blood is 1.5(mg/dL). The percentage of the blood leaving his heart at each contraction is 42. Will the patient be alive or dead?

Confusion matrix for Random forest - un-scaled

```

# Make predictions on the un-scaled test set
predictions <- predict(tree.Heart_disease_unscaled, newdata=testing_set_unscaled,
                        type="class")

confusionMatrix(testing_set_unscaled$DEATH_EVENT, predictions)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0  54   7
##           1   7  22
##
##               Accuracy : 0.8444
##               95% CI : (0.7528, 0.9123)
##       No Information Rate : 0.6778

```

```
##      P-Value [Acc > NIR] : 0.0002694
##
##              Kappa : 0.6439
##
## Mcnemar's Test P-Value : 1.0000000
##
##      Sensitivity : 0.8852
##      Specificity : 0.7586
##      Pos Pred Value : 0.8852
##      Neg Pred Value : 0.7586
##      Prevalence : 0.6778
##      Detection Rate : 0.6000
##      Detection Prevalence : 0.6778
##      Balanced Accuracy : 0.8219
##
##      'Positive' Class : 0
##
```

When the un-scaled data is fitted to decision tree model, we obtained the accuracy of 75.5%

MODEL 3 - DECISION TREES

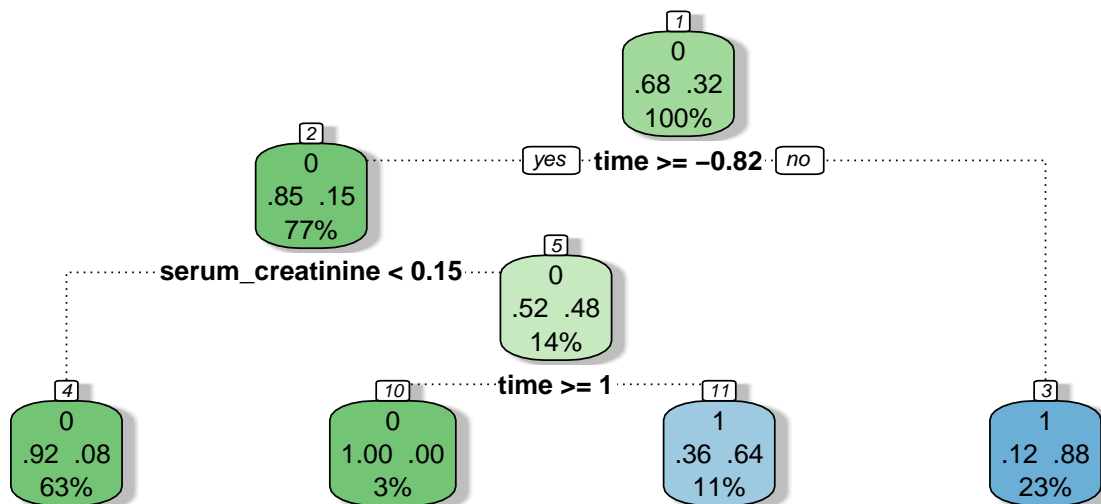
MODEL 3.b - DECISION TREES - ALL VARIABLES - SCALED

```
tree.Heart_disease_scaled = tree(DEATH_EVENT~.,training_set_scaled)
summary(tree.Heart_disease_scaled)
```

```
##
## Classification tree:
## tree(formula = DEATH_EVENT ~ ., data = training_set_scaled)
## Variables actually used in tree construction:
## [1] "time"                "serum_creatinine"
## [3] "sex"                 "ejection_fraction"
## [5] "age"                 "creatinine_phosphokinase"
## [7] "high_blood_pressure" "platelets"
## Number of terminal nodes: 14
## Residual mean deviance: 0.3906 = 76.16 / 195
## Misclassification error rate: 0.09569 = 20 / 209
```

Plotting the tree

```
#Better tree visualization using rpart package
tree.Heart_disease_scaled = rpart(DEATH_EVENT~., training_set_scaled)
fancyRpartPlot(tree.Heart_disease_scaled)
```



Rattle 2023-May-16 13:31:19 vishaypaka

Confusion matrix for Decision tree - scaled

```
# Make predictions on the scaled test set
predictions1 <- predict(tree.Heart_disease_scaled, newdata=testing_set_scaled,
                        type="class")
```

```
confusionMatrix(testing_set_scaled$DEATH_EVENT, predictions1)
```

Confusion Matrix and Statistics

##

Reference

Prediction 0 1

0 52 9

1 8 21

##

Accuracy : 0.8111

95% CI : (0.7149, 0.8859)

No Information Rate : 0.6667

P-Value [Acc > NIR] : 0.001794

##

Kappa : 0.5714

##

McNemar's Test P-Value : 1.000000

##

Sensitivity : 0.8667

Specificity : 0.7000

Pos Pred Value : 0.8525

```
##          Neg Pred Value : 0.7241
##          Prevalence : 0.6667
##          Detection Rate : 0.5778
##    Detection Prevalence : 0.6778
##          Balanced Accuracy : 0.7833
##
##          'Positive' Class : 0
##
```


When the scaled data is fitted to decision tree model, we obtained the accuracy of 81.1%. We can observe the Decision tree algorithm performs well on un-scaled data compared to scaled data. So, the model with un-scaled data is selected as best model in Decision tree with 84.4 accuracy and will compare it with other models later in this research.

MODEL 4 - RANDOM FOREST ALGORITHM

MODEL 4.a -> RANDOM FOREST ALGORITHM - ALL VARIABLES - UNSCALED

```
#tuning the mtry parameter using the tuneRF function from the randomForest package.
set.seed(100)
tune.rf <- tuneRF(training_set_unscaled[, -13], training_set_unscaled$DEATH_EVENT,
                  ntreeTry=100, stepFactor=1.5, plot=TRUE, dobest=TRUE)
```

```
## mtry = 3   OOB error = 19.14%
## Searching left ...
## mtry = 2   OOB error = 19.14%
## 0 0.05
## Searching right ...
## mtry = 4   OOB error = 17.22%
## 0.1 0.05
## mtry = 6   OOB error = 20.57%
## -0.1944444 0.05
```



The lower the out-of-bag error rate the best the model gives the accuracy. So, from above chart analysis we can set $m_{try} = 4$.

```
set.seed(0)
rf.Heart_disease_unscaled = randomForest(DEATH_EVENT~., data=training_set_unscaled, mtry=4,
                                         importance=TRUE)

predicted_values_rf = predict(rf.Heart_disease_unscaled, newdata = testing_set_unscaled)
```

Confusion matrix for Random forest - un-scaled

```
#Confusion Matrix - UNSCALED MODEL
confusionMatrix(predicted_values_rf, testing_set_unscaled$DEATH_EVENT)
```


```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 57  6
##           1  4 23
##
##           Accuracy : 0.8889
##           95% CI : (0.8051, 0.9454)
##           No Information Rate : 0.6778
```

```
##      P-Value [Acc > NIR] : 2.826e-06
##
##              Kappa : 0.7409
##
## Mcnemar's Test P-Value : 0.7518
##
##      Sensitivity : 0.9344
##      Specificity : 0.7931
##      Pos Pred Value : 0.9048
##      Neg Pred Value : 0.8519
##      Prevalence : 0.6778
##      Detection Rate : 0.6333
##      Detection Prevalence : 0.7000
##      Balanced Accuracy : 0.8638
##
##      'Positive' Class : 0
##
```

When the un-scaled data is fitted to random forest model, we obtained the accuracy of 88.89%

```
# Plot variable importance to determine which variables are most important.
varImpPlot(rf.Heart_disease_unscaled)
```

rf.Heart_disease_unscaled



MODEL 4.b -> RANDOM FOREST ALGORITHM - ALL VARIABLES - SCALED

```
set.seed(100)
rf.Heart_disease_scaled = randomForest(DEATH_EVENT~.,data=training_set_scaled, mtry=4,
                                       importance=TRUE)

predicted_values_rf1 = predict(rf.Heart_disease_scaled,newdata = testing_set_scaled)
```

Confusion matrix for Random forest - scaled


```
#Confusion Matrix - SCALED MODEL
confusionMatrix(predicted_values_rf1, testing_set_scaled$DEATH_EVENT)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0   1
##           0 50   8
##           1 11  21
##
##              Accuracy : 0.7889
##              95% CI : (0.6901, 0.8679)
##      No Information Rate : 0.6778
##      P-Value [Acc > NIR] : 0.01376
##
##              Kappa : 0.5294
##
##  Mcnemar's Test P-Value : 0.64636
##
##              Sensitivity : 0.8197
##              Specificity : 0.7241
##              Pos Pred Value : 0.8621
##              Neg Pred Value : 0.6562
##              Prevalence : 0.6778
##              Detection Rate : 0.5556
##      Detection Prevalence : 0.6444
##      Balanced Accuracy : 0.7719
##
##              'Positive' Class : 0
##
```

When the scaled data is fitted to random forest model, we obtained the accuracy of 78.89% We can observe the random forest algorithm performs well on un-scaled data compared to scaled data. So, the model with un-scaled data is selected as best model in random forest with 88.89 accuracy and will compare it with other models below.

```
# Plot variable importance to determine which variables are most important.
varImpPlot(rf.Heart_disease_scaled)
```

rf.Heart_disease_scaled



By comparing both the plots of importance variables from scaled and un-scaled data, we can tell that the predictors : time, serum creatinine, serum sodium,ejection fraction and age are the most important features for the prediction and are most predictive.

COMPARING ALL THE MODELS.

```
#Collecting all the best accuracy and sensitivity(recall) metrics from out all the models(scaled/unscal
acc_lr = 88.89
acc_knn = 85.56
acc_dtree = 84.44
acc_rf = 88.89
tpr_lr = 93.44
tpr_knn = 87.50
tpr_dtree = 88.52
tpr_rf = 93.44
str(Heart_disease)
```

```
## 'data.frame': 299 obs. of 13 variables:
## $ age : num 75 55 65 50 65 90 75 60 65 80 ...
## $ anaemia : Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 2 1 2 ...
## $ creatinine_phosphokinase: num 582 7861 146 111 160 ...
## $ diabetes : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 2 1 1 ...
## $ ejection_fraction : num 20 38 20 20 20 40 15 60 65 35 ...
## $ high_blood_pressure : Factor w/ 2 levels "0","1": 2 1 1 1 1 2 1 1 1 2 ...
```

```
## $ platelets           : num  265000 263358 162000 210000 327000 ...
## $ serum_creatinine    : num   1.9 1.1 1.3 1.9 2.7 2.1 1.2 1.1 1.5 9.4 ...
## $ serum_sodium        : num   20 38 20 20 20 40 15 60 65 35 ...
## $ sex                 : Factor w/ 2 levels "0","1": 2 2 2 2 1 2 2 2 1 2 ...
## $ smoking             : Factor w/ 2 levels "0","1": 1 1 2 1 1 2 1 2 1 2 ...
## $ time                : int    4 6 7 7 8 8 10 10 10 10 ...
## $ DEATH_EVENT         : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
```

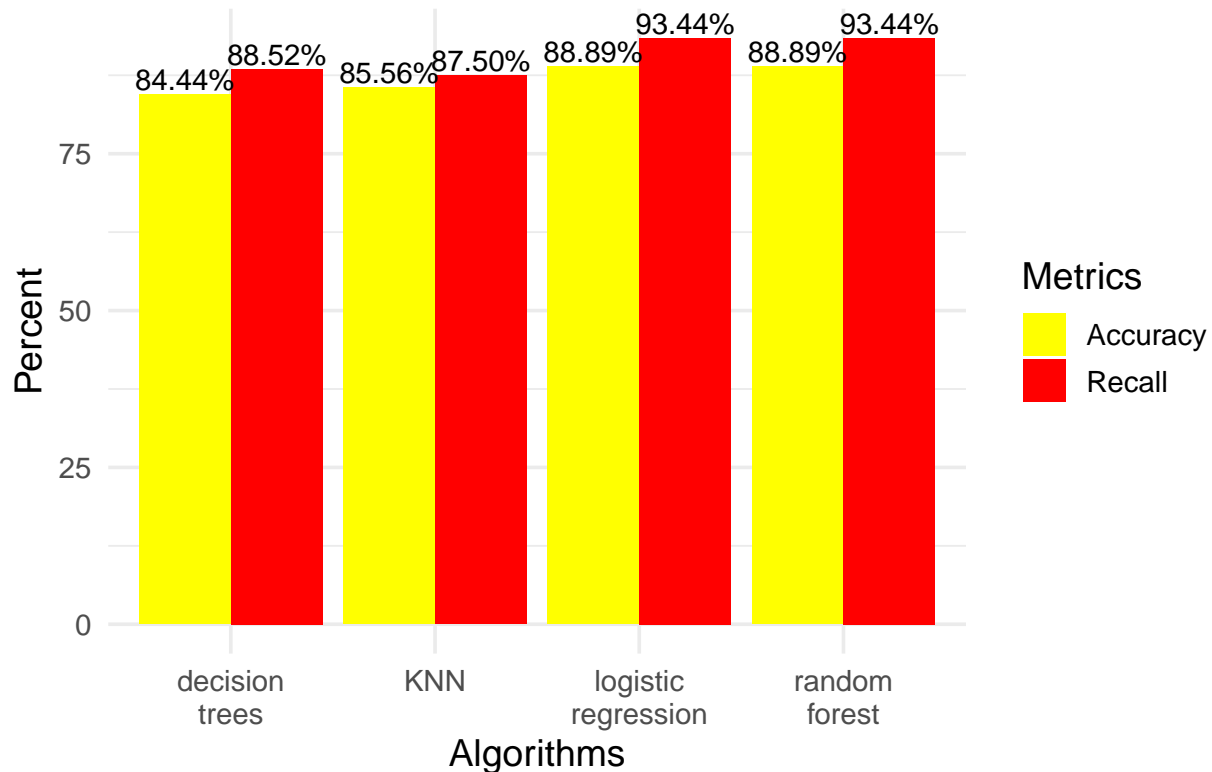
PLOTTING THE METRICS OF ALL DIFFERENT CLASSIFICATION MODELS(BEST ONE)

```
# create a data frame with algorithm names, accuracy, and recall
df <- tibble(
  Algorithm = c("logistic\nregression", "KNN", "decision\ntrees", "random\nforest"),
  Accuracy = c(acc_lr, acc_knn, acc_dtree, acc_rf),
  Recall = c(tpr_lr, tpr_knn, tpr_dtree, tpr_rf)
)

# reshape the data frame into a longer format
df_long <- df %>%
  pivot_longer(cols = -Algorithm, names_to = "Metrics", values_to = "Percent")

# plot the data using ggplot2
ggplot(df_long, aes(x = reorder(Algorithm, X = Percent),
  y = Percent,
  fill = Metrics)) +
  geom_bar(stat = "identity",
    position = "dodge",
    alpha=1.0) +
  geom_text(aes(group = Metrics, label = str_c(sprintf("%.2f", Percent), "%")),
    position = position_dodge(width = 0.9), vjust = -0.2) +
  scale_fill_manual(values = c("yellow", "red")) +
  labs(x = "Algorithms", title = "Metrics of different classification models performed") +
  theme_minimal(base_size = 14)
```

Metrics of different classification models performed



We can observe the accuracy and recall(sensitivity) of all the models which are plotted above the bars.

From the above plot we can say that logistic regression is the “best” model in terms of recall, while KNN is the “best” model in terms of accuracy.

Now let’s try to do simulation and add some data.

SIMULATION OF THE DATA BASED ON EXISTING DATA

Using random normal distributions with mean and standard deviation values taken from the original Heart_disease dataset. we will create a new dataset of 1000 observations with the same columns as the original, but with different values generated by the random distributions. The new dataset can be used for testing or training machine learning models without having to collect new data. We are creating 1000 observations because the simulated dataset should not impact more than the original dataset.

```
set.seed(1400)

N = 1400
simulated_heart_disease = data.frame(
  age = round(rnorm(N, mean(Heart_disease$age), sd(Heart_disease$age))),
  anaemia = sample(c(0, 1), N, replace = TRUE),
  creatinine_phosphokinase = round(rnorm(N, mean(Heart_disease$creatinine_phosphokinase),
                                          sd(Heart_disease$creatinine_phosphokinase))),
  diabetes = sample(c(0, 1), N, replace = TRUE),
  ejection_fraction = round(rnorm(N, mean(Heart_disease$ejection_fraction),
                                   sd(Heart_disease$ejection_fraction))),
```

```

high_blood_pressure = sample(c(0, 1), N, replace = TRUE),
platelets = round(rnorm(N, mean(Heart_disease$platelets), sd(Heart_disease$platelets))),
serum_creatinine = round(rnorm(N, mean(Heart_disease$serum_creatinine),
                                sd(Heart_disease$serum_creatinine))),
serum_sodium = round(rnorm(N, mean(Heart_disease$serum_sodium),
                                sd(Heart_disease$serum_sodium))),
sex = sample(c(0, 1), N, replace = TRUE),
smoking = sample(c(0, 1), N, replace = TRUE),
time = round(rnorm(N, mean(Heart_disease$time), sd(Heart_disease$time)))
)

simulated_heart_disease$anaemia <- factor(simulated_heart_disease$anaemia, levels = c(0,1))
simulated_heart_disease$diabetes <- factor(simulated_heart_disease$diabetes, levels = c(0,1))
simulated_heart_disease$high_blood_pressure <- factor(simulated_heart_disease$high_blood_pressure,
                                                    levels = c(0,1))
simulated_heart_disease$sex <- factor(simulated_heart_disease$sex, levels = c(0,1))
simulated_heart_disease$smoking <- factor(simulated_heart_disease$smoking, levels = c(0,1))

#Fitting the LOGISTIC MODEL
#predicting on the un-scaled testing data
simulated_prediction <- predict(rf.Heart_disease_unscaled, newdata=simulated_heart_disease)
simulated_data <- cbind(simulated_heart_disease, simulated_prediction)

```

Combing heart disease dataset with the obtained simulation dataset

```

colnames(simulated_data)[13] <- "DEATH_EVENT"
heart_disease_simulated = rbind(training_set_unscaled, testing_set_unscaled, simulated_data)

```

SPLITTING INTO TRAINING AND TESTING

```

set.seed(1001)
train_simulated <- sample(nrow(heart_disease_simulated), 0.8 * nrow(heart_disease_simulated))
training_set_simulated <- heart_disease_simulated[train_simulated, ]
testing_set_simulated <- heart_disease_simulated[-train_simulated, ]

```

FITTING THE RANDOM FOREST MODEL AND PREDICTING RESULTS FOR SIMULATED DATA

```

#fitting the un-scaled data on random forest model
set.seed(100)
rf.Heart_disease_simulated_unscaled = randomForest(DEATH_EVENT~., data=training_set_simulated,
                                                    mtry=6, importance=TRUE)

predict_test_all_simulated = predict(rf.Heart_disease_simulated_unscaled,
                                      newdata = testing_set_simulated)

```

Confusion matrix for full simulated model on random forest.

```
confusionMatrix(predict_test_all_simulated,testing_set_simulated$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 254  13
##           1   8  65
##
##           Accuracy : 0.9382
##           95% CI : (0.9071, 0.9614)
##       No Information Rate : 0.7706
##       P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.8213
##
##  McNemar's Test P-Value : 0.3827
##
##           Sensitivity : 0.9695
##           Specificity : 0.8333
##       Pos Pred Value : 0.9513
##       Neg Pred Value : 0.8904
##           Prevalence : 0.7706
##       Detection Rate : 0.7471
##       Detection Prevalence : 0.7853
##       Balanced Accuracy : 0.9014
##
##       'Positive' Class : 0
##
```

From the added simulated data to our original dataset we got an accuracy of 93.82% and an sensitivity (recall) of 96.95%. Which is very good than all the models we performed above.

FITTING THE LOGISTIC MODEL AND PREDICTING RESULTS

```
#fitting the un-scaled data on logistic regression model
glm.all_simulated = glm(DEATH_EVENT~.,data=training_set_simulated,family="binomial")

#predicting on the testing un-scaled data
predict_test_all_simulated1 <- factor(ifelse(predict(glm.all_simulated,
                                                    testing_set_simulated ,
                                                    type ="response") > 0.5, "1", "0"))

summary(glm.all_simulated)

##
## Call:
## glm(formula = DEATH_EVENT ~ ., family = "binomial", data = training_set_simulated)
##
```

```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4635  -0.4706  -0.1486   0.1656   3.4745
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    7.467e-01  7.116e-01   1.049   0.2940
## age            5.613e-02  8.283e-03   6.776 1.23e-11 ***
## anaemia1       4.378e-01  1.811e-01   2.418   0.0156 *
## creatinine_phosphokinase 4.395e-04  9.849e-05   4.463 8.10e-06 ***
## diabetes1     -1.795e-01  1.803e-01  -0.996   0.3194
## ejection_fraction -3.168e-02  7.907e-03  -4.006 6.17e-05 ***
## high_blood_pressure1 3.541e-01  1.805e-01   1.962   0.0498 *
## platelets     -4.996e-07  9.071e-07  -0.551   0.5818
## serum_creatinine 6.232e-01  9.241e-02   6.744 1.54e-11 ***
## serum_sodium   -4.654e-02  7.991e-03  -5.824 5.73e-09 ***
## sex1          -4.090e-01  1.811e-01  -2.259   0.0239 *
## smoking1       1.819e-01  1.798e-01   1.011   0.3119
## time          -3.358e-02  2.092e-03 -16.050 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1579.44  on 1358  degrees of freedom
## Residual deviance:  802.85  on 1346  degrees of freedom
## AIC: 828.85
##
## Number of Fisher Scoring iterations: 6
```

For the added simulated data to original dataset, we can see there are added significant variables to our full model, which are age, sex, ejection fraction, serum creatinine, time and, creatinine phosphokinase.

Confusion matrix for full simulated model

```
confusionMatrix(predict_test_all_simulated1,testing_set_simulated$DEATH_EVENT)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##      0 245  21
##      1  17  57
##
##              Accuracy : 0.8882
##              95% CI : (0.8498, 0.9197)
##      No Information Rate : 0.7706
##      P-Value [Acc > NIR] : 1.994e-08
##
##              Kappa : 0.6781
##
##      McNemar's Test P-Value : 0.6265
```

```
##
##          Sensitivity : 0.9351
##          Specificity : 0.7308
##          Pos Pred Value : 0.9211
##          Neg Pred Value : 0.7703
##          Prevalence : 0.7706
##          Detection Rate : 0.7206
##          Detection Prevalence : 0.7824
##          Balanced Accuracy : 0.8329
##
##          'Positive' Class : 0
##
```

From the added simulated data to our original dataset we got an accuracy of 88.82% and an sensitivity (recall) of 93.51% when fitted with logistic model.

So, the Random forest model fitted with simulated data is very good than logistic model we performed above with an accuracy of 93.82% when fitted with random forest algorithm.

Tabulating the results of all models.

```
Classification_models = c('Logistic Model - All Variables(unscaled)',
                          'Logistic Model - All Variables(scaled)',
                          'Logistic Model - Removing serum_sodium Variable(unscaled)',
                          'Logistic Model - Removing serum_sodium Variable(scaled)',
                          'Logistic Model - Significant Variables(unscaled)',
                          'Logistic Model - Significant Variables(scaled)',
                          'KNN Model - All continuous variables',
                          'KNN Model - significant Continuous variables',
                          'Decision tree Model (unscaled)', 'Decision tree Model (scaled)',
                          'Random-forest Model (unscaled)', 'Random-forest Model (scaled)',
                          'Random-forest Model - All Variables (After adding simulated data)',
                          'Logistic Model - All Variables (After adding simulated data)')
Accuracy_rates = c('88.89%', '83.33%', '88.89%', '83.33%', '88.89%',
                   '82.22%', '61.11%', '85.56%', '84.44%', '81.11%', '88.89%', '78.89%', '93.82%', '88.82%')

All_models_comparision = data.frame(Classification_models, Accuracy_rates)
All_models_comparision
```

```
##          Classification_models
## 1      Logistic Model - All Variables(unscaled)
## 2      Logistic Model - All Variables(scaled)
## 3      Logistic Model - Removing serum_sodium Variable(unscaled)
## 4      Logistic Model - Removing serum_sodium Variable(scaled)
## 5      Logistic Model - Significant Variables(unscaled)
## 6      Logistic Model - Significant Variables(scaled)
## 7      KNN Model - All continuous variables
## 8      KNN Model - significant Continuous variables
## 9      Decision tree Model (unscaled)
## 10     Decision tree Model (scaled)
## 11     Random-forest Model (unscaled)
## 12     Random-forest Model (scaled)
```



```
## 13 Random-forest Model - All Variables (After addidng simulated data)
## 14      Logistic Model - All Variables (After addidng simulated data)
##      Accuracy_rates
## 1      88.89%
## 2      83.33%
## 3      88.89%
## 4      83.33%
## 5      88.89%
## 6      82.22%
## 7      61.11%
## 8      85.56%
## 9      84.44%
## 10     81.11%
## 11     88.89%%
## 12     78.89%
## 13     93.82%
## 14     88.82%
```

Performing chi-square test.

```
# Create a contingency table between smoking and death event
smoking_death_table <- table(Heart_disease$smoking, Heart_disease$DEATH_EVENT)
```

```
# Print the contingency table
smoking_death_table
```

```
##
##      0    1
## 0 137  66
## 1   66  30
```

```
# Perform the chi-square test
chisq.test(smoking_death_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  smoking_death_table
## X-squared = 0.0073315, df = 1, p-value = 0.9318
```

```
# Create a contingency table between sex and death event
sex_death_table <- table(Heart_disease$sex, Heart_disease$DEATH_EVENT)
```

```
# Print the contingency table
sex_death_table
```

```
##
##      0    1
## 0   71  34
## 1  132  62
```

```
# Perform the chi-square test
chisq.test(sex_death_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: sex_death_table
## X-squared = 0, df = 1, p-value = 1
```

```
# Create a contingency table between high blood pressure and death event
high_blood_pressure_death_table <- table(Heart_disease$high_blood_pressure,
                                          Heart_disease$DEATH_EVENT)
```

```
# Print the contingency table
high_blood_pressure_death_table
```

```
##
##      0    1
## 0 137  57
## 1   66  39
```

```
# Perform the chi-square test
chisq.test(high_blood_pressure_death_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: high_blood_pressure_death_table
## X-squared = 1.5435, df = 1, p-value = 0.2141
```

```
# Create a contingency table between diabetes and death event
diabetes_death_table <- table(Heart_disease$diabetes, Heart_disease$DEATH_EVENT)
```

```
# Print the contingency table
diabetes_death_table
```

```
##
##      0    1
## 0 118  56
## 1   85  40
```

```
# Perform the chi-square test
chisq.test(diabetes_death_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: diabetes_death_table
## X-squared = 2.1617e-30, df = 1, p-value = 1
```

```
# Create a contingency table between anaemia and death event
anaemia_death_table <- table(Heart_disease$anaemia, Heart_disease$DEATH_EVENT)
```

```
# Print the contingency table
anaemia_death_table
```

```
##
##      0    1
## 0 120  50
## 1   83  46
```

```
# Perform the chi-square test
chisq.test(anaemia_death_table)
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  anaemia_death_table
## X-squared = 1.0422, df = 1, p-value = 0.3073
```

Analysing the Chi-square results with death event and making conclusion

```
Predictor_variable = c("anaemia","diabetes","High blood pressure","sex","smoking")
xsquared = c('0.0073315','0','1.5435','2.1617e-30','1.0422')
pvalue = c('0.9318','1','0.2141','1','0.3073')
```

```
Chi_square_df = data.frame(Predictor_variable,xsquared,pvalue)
```

```
Chi_square_df
```

```
##   Predictor_variable  xsquared pvalue
## 1          anaemia 0.0073315 0.9318
## 2          diabetes         0      1
## 3 High blood pressure    1.5435 0.2141
## 4              sex 2.1617e-30      1
## 5           smoking    1.0422 0.3073
```

The above table shows the results of chi-squared tests performed to test the association between various predictor variables (anaemia, diabetes, high blood pressure, sex, smoking) and the death event in the heart disease dataset.

The x-squared value indicates the degree of association between the predictor variable and the outcome variable (death event). A higher value of x-squared indicates a stronger association between the predictor variable and the outcome variable.

Based on the above table, we can conclude that the smoking and diabetes are weakly associated with the death event, as indicated by their relatively low x-squared values and high p-values.

Performing multiple linear regression on time(follow-up days)

To predict the effect of age, ejection fraction, serum creatinine level, and other clinical features on the patient's follow-up period or time to death:

```
# Fit the multiple linear regression model
model <- lm(time ~ age + ejection_fraction + serum_creatinine + anaemia + diabetes +
             high_blood_pressure + sex + smoking +
             +platelets + creatinine_phosphokinase, data = Heart_disease)

# Print the summary of the model
summary(model)
```

```
##
## Call:
## lm(formula = time ~ age + ejection_fraction + serum_creatinine +
##     anaemia + diabetes + high_blood_pressure + sex + smoking +
##     +platelets + creatinine_phosphokinase, data = Heart_disease)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -154.204  -58.029   -6.694   66.368  154.128
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.272e+02  3.062e+01   7.420 1.32e-12 ***
## age           -1.204e+00  3.732e-01  -3.225  0.00141 **
## ejection_fraction  3.558e-01  3.692e-01   0.964  0.33600
## serum_creatinine -8.681e+00  4.219e+00  -2.057  0.04056 *
## anaemia1       -2.040e+01  8.929e+00  -2.285  0.02306 *
## diabetes1      -1.673e-01  8.925e+00  -0.019  0.98506
## high_blood_pressure1 -2.971e+01  9.114e+00  -3.260  0.00125 **
## sex1           -5.096e-01  1.040e+01  -0.049  0.96094
## smoking1       -6.917e+00  1.041e+01  -0.664  0.50694
## platelets      -1.585e-06  4.489e-05  -0.035  0.97186
## creatinine_phosphokinase -4.905e-03  4.550e-03  -1.078  0.28192
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 74.2 on 288 degrees of freedom
## Multiple R-squared:  0.1168, Adjusted R-squared:  0.0861
## F-statistic: 3.807 on 10 and 288 DF,  p-value: 7.627e-05
```

We can observe from the above summary that age, serum creatinine, anaemia, high blood pressure are statistically significant predictors of the dependent variable. This suggests that older patients, patients with higher serum creatinine levels, anaemic patients, and patients with high blood pressure are more likely to have shorter follow-up periods or higher risk of death. But the adjusted R-squared value is only 0.08615 suggests that the model explains only 8.6% of the variance in the dependent variable. The low adjusted R-squared value suggests that the model may not be a good fit for the data and may not provide accurate predictions.